

Variances add when independent

$$\begin{aligned} \text{Var}(X + Y) &= (E(X + Y)^2) - E((X + Y)^2) \\ &= (E(X) + E(Y))^2 - E(X^2 + 2XY + Y^2) \\ &= E(X)^2 + E(Y)^2 + 2E(X)E(Y) - E(X^2) - 2E(XY) - E(Y^2) \\ &= \text{Var}(X) + \text{Var}(Y) + 2(E(X)E(Y) - E(XY)) \end{aligned}$$

So it remains to show that the quantity $E(X)E(Y) - E(XY) = 0$ when X and Y are independent. In fact, this quantity, which is easily shown to be equal to $E((X - E(X))(Y - E(Y)))$ by expanding, is called Covariance ($\text{Cov}(X, Y)$).

When X and Y are independent, then for each possible value of X , the expected value of $Y - E(Y) = 0$. Since the expectation when there are a bunch of possible outcomes is intuitively equal to the sum of the expected value from those outcomes times the probability of those outcomes, we can apply this to the possible values of X , so the quantity $E((X - E(X))(Y - E(Y)))$ becomes a weighted sum of terms like $E((x - E(X))P(X=x)(Y - E(Y)))$, which is a bunch of terms like $E(\text{constant}(Y - E(Y)))$ which is 0 because $Y - E(Y)$ is always 0 regardless of the value of X due to the fact that X and Y are independent.

An important idea is that independence allows you to take $E(X)E(Y) = E(XY)$.

Normal distribution formula and variance

Formula:

So we have a probability distribution based on e^{-x^2} , this is an assumption. We change this to $e^{-\frac{x^2}{2}}$ since it turns out (we will prove this after) that this will make the variance equal 1. Then we have to rescale by a factor, which turns out to be $\sqrt{2\pi}$ (This is what we will prove, but really we just need to know that the area exists which is easy to prove, however the square root of pi is so fun that I can't just not show the proof) so that the area under the curve equals 1. We will prove that the total area under e^{-x^2} is $\sqrt{\pi}$ then it is clear that by shifting the mean the area will not change and that if we divide x by a constant σ then we are stretching the curve and therefore its area by a factor of σ so we need to divide the correction factor by σ to compensate and keep the area under the curve at 1.

Now, we will consider what happens if we take the e^{-x^2} curve and rotate it about the y axis, like this:




Image: 2 dimensional normal plotted on 3d graph

As you can see, it now essentially becomes a function of the distance from the origin d , ie $z = e^{-d^2}$. Now by pythagoras this becomes $z = e^{-(x^2+y^2)}$. We want to show that the volume of this surface is π ,

because then we know that if the value $\int_{-\infty}^{\infty} e^{-x^2} dx$ is k then the value of, the strip of this surface for a certain y value with a small width dy (As illustrated in the figure below) is approximately


$$\left(\int_{x=-\infty}^{\infty} e^{-(y^2+x^2)} dx\right) dy = (e^{-y^2} \int_{-\infty}^{\infty} e^{-x^2} dx) dy = (e^{-y^2} k) dy,$$


Image: A thin slice of the 2d normal in the above image

Then the volume under the surface is the sum of the volumes of these tiny strips (as dy gets smaller), ie, as we have seen in the interpretation of the integral as a sum, $\int_{-\infty}^{\infty} (e^{-y^2} k) dy = k^2$. Therefore, if we can show that the volume under the surface which is k^2 is π , then $k = \sqrt{\pi}$ as required.

The normal distribution is very special in the sense that if we take x and y independently normally distributed (which makes the probability density become the product of the probability density for x and y), the resulting distribution has rotational symmetry and depends only on the instance, as it becomes $e^{-(x^2+y^2)}$, which is a function of the distance by pythagoras.

Now, to prove that the volume is indeed π , we will consider what happens if we split the volume into concentric rings, like this:




Image: 2d normal on 3d graph approximated by concentric rings of the right height

Suppose the width of the rings is dr and their distance from the origin is r , then the height of the rings is e^{-r^2} and the circumference of them is $2\pi r$ so the volume of each ring is $2\pi r e^{-r^2} dr$. So, the total volume under the figure is $\int_0^{\infty} 2\pi r e^{-r^2} dr$, again by the interpretation of integration as a sum. Note that you can check by differentiating that an antiderivative of $2\pi r e^{-r^2}$ is $-\pi e^{-r^2}$. So, we have to evaluate $[-\pi e^{-r^2}]_0^{\infty}$ which is $0 - [-\pi]$ which is π as required.

Comment: This is often the case in mathematics, the idea seeing things in 2 different ways is extremely common (like here with strips vs rings).

Variance:

If we have a standard normal $\frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\infty} e^{-\frac{x^2}{2}} dx$ then the variance is given by $E(x^2)$ since the mean is 0 by symmetry so we need to work out $\frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\infty} x^2 e^{-\frac{x^2}{2}} dx$ and show that it is indeed 1, then it becomes clear that the variance is always σ^2 due to scaling properties of variances (ie, if we were to stretch the normal distribution by $2x$, we get this effect by multiplying σ by 2 in the formula and the variance multiplies by 4 by variance scaling properties and the variance is unchanged by shift in the mean since it is related to distance from mean)

Now, we try integrating $\frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\infty} e^{-\frac{x^2}{2}} dx$ (which, keep in mind, is 1) by parts with $e^{-\frac{x^2}{2}}$ being the part to differentiate and 1 being the part to integrate. This will give us $\frac{1}{\sqrt{2\pi}} \left[\left[x e^{-\frac{x^2}{2}} \right]_{-\infty}^{\infty} + \int_{-\infty}^{\infty} x^2 e^{-\frac{x^2}{2}} dx \right]$ (the – in the integration by parts formula cancels with the – from differentiating the exponential). Now, since $x e^{-\frac{x^2}{2}}$ goes to 0 as x goes to positive and negative infinity (This is because clearly the derivative of a normal distribution must go to 0 on either side and this is just minus that, or alternatively because the exponential term decays much faster than x grows), this formula, which is equal to 1, reduces to the integral $\frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\infty} x^2 e^{-\frac{x^2}{2}} dx$ that we needed to find.

Regression line formula minimizes squares

Theorem: Given a [simple linear regression model](#) with independent observations

$$y = \beta_0 + \beta_1 x + \varepsilon, \varepsilon_i \sim \mathcal{N}(0, \sigma^2), i = 1, \dots, n, \quad (1)$$

the parameters minimizing the [residual sum of squares](#) are given by

$$\begin{aligned} \hat{\beta}_0 &= \bar{y} - \hat{\beta}_1 \bar{x} \\ \hat{\beta}_1 &= \frac{s_{xy}}{s_x^2} \end{aligned} \quad (2)$$

where $\bar{x}$ and $\bar{y}$ are the [sample means](#), s_x^2 is the [sample variance](#) of x and s_{xy} is the [sample covariance](#) between x and y .

Proof: The [residual sum of squares](#) is defined as

$$\text{RSS}(\beta_0, \beta_1) = \sum_{i=1}^n \varepsilon_i^2 = \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2. \quad (3)$$

The derivatives of $\text{RSS}(\beta_0, \beta_1)$ with respect to β_0 and β_1 are

$$\begin{aligned} \frac{d\text{RSS}(\beta_0, \beta_1)}{d\beta_0} &= \sum_{i=1}^n 2(y_i - \beta_0 - \beta_1 x_i)(-1) \\ &= -2 \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i) \\ \frac{d\text{RSS}(\beta_0, \beta_1)}{d\beta_1} &= \sum_{i=1}^n 2(y_i - \beta_0 - \beta_1 x_i)(-x_i) \end{aligned} \quad (4)$$

$$= -2 \sum_{i=1}^n (x_i y_i - \beta_0 x_i - \beta_1 x_i^2)$$

and setting these derivatives to zero

$$\begin{aligned} 0 &= -2 \sum_{i=1}^n (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i) \\ 0 &= -2 \sum_{i=1}^n (x_i y_i - \hat{\beta}_0 x_i - \hat{\beta}_1 x_i^2) \end{aligned} \quad (5)$$

yields the following equations:

$$\begin{aligned} \hat{\beta}_1 \sum_{i=1}^n x_i + \hat{\beta}_0 \cdot n &= \sum_{i=1}^n y_i \\ \hat{\beta}_1 \sum_{i=1}^n x_i^2 + \hat{\beta}_0 \sum_{i=1}^n x_i &= \sum_{i=1}^n x_i y_i. \end{aligned} \quad (6)$$

From the first equation, we can derive the estimate for the intercept:

$$\begin{aligned} \hat{\beta}_0 &= \frac{1}{n} \sum_{i=1}^n y_i - \hat{\beta}_1 \cdot \frac{1}{n} \sum_{i=1}^n x_i \\ &= \bar{y} - \hat{\beta}_1 \bar{x}. \end{aligned} \quad (7)$$

From the second equation, we can derive the estimate for the slope:

$$\begin{aligned} \hat{\beta}_1 \sum_{i=1}^n x_i^2 + \hat{\beta}_0 \sum_{i=1}^n x_i &= \sum_{i=1}^n x_i y_i \\ \hat{\beta}_1 \sum_{i=1}^n x_i^2 + (\bar{y} - \hat{\beta}_1 \bar{x}) \sum_{i=1}^n x_i &\stackrel{(7)}{=} \sum_{i=1}^n x_i y_i \\ \hat{\beta}_1 \left(\sum_{i=1}^n x_i^2 - \bar{x} \sum_{i=1}^n x_i \right) &= \sum_{i=1}^n x_i y_i - \bar{y} \sum_{i=1}^n x_i \\ \hat{\beta}_1 &= \frac{\sum_{i=1}^n x_i y_i - \bar{y} \sum_{i=1}^n x_i}{\sum_{i=1}^n x_i^2 - \bar{x} \sum_{i=1}^n x_i}. \end{aligned} \quad (8)$$

Note that the numerator can be rewritten as

$$\begin{aligned} \sum_{i=1}^n x_i y_i - \bar{y} \sum_{i=1}^n x_i &= \sum_{i=1}^n x_i y_i - n \bar{x} \bar{y} \\ &= \sum_{i=1}^n x_i y_i - n \bar{x} \bar{y} - n \bar{x} \bar{y} + n \bar{x} \bar{y} \\ &= \sum_{i=1}^n x_i y_i - \bar{y} \sum_{i=1}^n x_i - \bar{x} \sum_{i=1}^n y_i + \sum_{i=1}^n \bar{x} \bar{y} \\ &= \sum_{i=1}^n (x_i y_i - x_i \bar{y} - \bar{x} y_i + \bar{x} \bar{y}) \\ &= \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}) \end{aligned} \quad (9)$$

and that the denominator can be rewritten as

$$\begin{aligned}
 \sum_{i=1}^n x_i^2 - \bar{x} \sum_{i=1}^n x_i &= \sum_{i=1}^n x_i^2 - n\bar{x}^2 \\
 &= \sum_{i=1}^n x_i^2 - 2n\bar{x}\bar{x} + n\bar{x}^2 \\
 &= \sum_{i=1}^n x_i^2 - 2\bar{x} \sum_{i=1}^n x_i + \sum_{i=1}^n \bar{x}^2 \\
 &= \sum_{i=1}^n (x_i^2 - 2\bar{x}x_i + \bar{x}^2) \\
 &= \sum_{i=1}^n (x_i - \bar{x})^2.
 \end{aligned} \tag{10}$$

With (9) and (10), the estimate from (8) can be simplified as follows:

$$\begin{aligned}
 \hat{\beta}_1 &= \frac{\sum_{i=1}^n x_i y_i - \bar{y} \sum_{i=1}^n x_i}{\sum_{i=1}^n x_i^2 - \bar{x} \sum_{i=1}^n x_i} \\
 &= \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2} \\
 &= \frac{\frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2} \\
 &= \frac{s_{xy}}{s_x^2}.
 \end{aligned} \tag{11}$$

Together, (7) and (11) constitute the ordinary least squares parameter estimates for simple linear regression.

PMCC between -1 and 1

Let $\vec{a} = (x_1 - \bar{x}, x_2 - \bar{x}, \dots, x_n - \bar{x})$ and $\vec{b} = (y_1 - \bar{y}, \dots, y_n - \bar{y})$. Then your formula is just $\frac{\vec{a} \cdot \vec{b}}{\|\vec{a}\| \|\vec{b}\|}$. Since $\vec{a} \cdot \vec{b} = \|\vec{a}\| \|\vec{b}\| \cos(\theta)$ your formula reduces to $\cos(\theta)$, where θ is the angle between $\vec{a}$ and $\vec{b}$ in n -dimensional euclidean space.

Noting that $\cos(\theta)$ is between -1 and 1 for real θ

PMCC measures correlation

I mean, if you look at the proof above, you see that it's literally the high dimensional cosine of the angle between your vector of x deviations from the mean and y deviations from the mean! It should then be obvious that if the deviations are close together, the cosine of the angle between them should be high, and vice versa, and furthermore the property should hold that scaling both of them shouldn't

change the angle between them, which takes care of that part. It also follows that PMCC is only -1 or 1 when x against y makes a straight line, as the deviation from the sample means of x and y in that case will always be the same, just scaled by a constant, thus the cosine of the angle between them will be -1 or 1.

Rank correlation coefficient formula

We compute the PMCC between the rank numbers.

x_i and y_i are integers 1 to n in any order

$$r = \frac{S_{xy}}{\sqrt{S_{xx}S_{yy}}} \quad \bar{y} = \frac{n+1}{2}$$

$$\bar{x} = \frac{n+1}{2}$$

S_{xy} - Will come back to this

$$S_{xx} = \sum (x_i - \bar{x})^2 = \sum_{i=1}^n \left(i - \frac{n+1}{2}\right)^2$$

$$= \sum_{i=1}^n i^2 - i(n+1) + \frac{(n+1)^2}{4}$$

Images: Shows my derivation on paper

$$= \sum_{i=1}^n i^2 - (n+1) \sum_{i=1}^n i + \sum_{i=1}^n \frac{(n+1)^2}{4}$$

$$= \frac{1}{6} n(n+1)(2n+1) - (n+1) \frac{n(n+1)}{2} + \frac{n(n+1)^2}{4}$$

$$= n(n+1) \left[\frac{1}{6} (2n+1) - \frac{n+1}{2} + \frac{n+1}{4} \right]$$

$$= n(n+1) \left[\frac{1}{3} n + \frac{1}{6} - \frac{1}{2} n - \frac{1}{2} + \frac{1}{4} n + \frac{1}{4} \right]$$

$$= n(n+1) \left[\frac{1}{12} n - \frac{1}{12} \right] = \frac{1}{12} n(n^2 - 1)$$

$$= \sum_{i=1}^n \left(x_i^r - \frac{5}{n+1}\right) \left(x_i^r - \frac{5}{n+1}\right)$$

$$z^{xx} = \sum_{i=1}^n (x_i^r - \bar{x})(x_i^r - \bar{x})$$

$$z = \frac{z^{xx} z^{xx}}{z^{xx} z^{xx}} = \frac{1(z^{xx})^2}{z^{xx} z^{xx}} = \frac{\frac{15 \cdot 16(15-1)}{1}}{15 \cdot 16} = \frac{15 \cdot 16(15-1)}{15 \cdot 16}$$

for integers 1 to n for x and y are both

$$\sum (x_i^r - \bar{x}) = \sum (x_i^r - 1)$$

~~$z^{xx} = z^{xx}$~~ ~~$z^{xx} = z^{xx}$~~ ~~$z^{xx} = z^{xx}$~~

$$\begin{aligned}
&= \sum_{i=1}^n x_i y_i - \frac{n+1}{2} x_i - \frac{n+1}{2} y_i + \frac{(n+1)^2}{4} \\
&= \sum_{i=1}^n x_i y_i - \frac{n+1}{2} \sum_{i=1}^n x_i - \frac{n+1}{2} \sum_{i=1}^n y_i + \sum_{i=1}^n \frac{(n+1)^2}{4} \\
&\quad \text{These terms are the sum from 1 to } n \text{ which is } \frac{n(n+1)}{2} \\
&= \sum_{i=1}^n x_i y_i - \frac{n(n+1)^2}{4} - \frac{n(n+1)^2}{4} + \frac{n(n+1)^2}{4} \\
&= \sum_{i=1}^n x_i y_i - \frac{n(n+1)^2}{4} = \sum_{i=1}^n x_i (x_i + y_i)
\end{aligned}$$

Note: $\sum_{i=1}^n x_i^2$ and $\sum_{i=1}^n y_i^2$ are the same as $\frac{1}{6}n(n+1)(2n+1)$ as they are adding the squares of the integers from 1 to n as those are the x_i and y_i values.

$$\begin{aligned}
S_{xy} &= \sum_{i=1}^n x_i y_i - \frac{1}{4}n^3 - \frac{1}{2}n^2 - \frac{1}{4}n \\
&= \sum_{i=1}^n x_i y_i + \frac{1}{12}n^3 - \frac{1}{12}n - \frac{1}{3}n^3 - \frac{1}{2}n^2 - \frac{1}{6}n
\end{aligned}$$

$$\begin{aligned}
&= \sum_{i=1}^n x_i y_i + \frac{1}{12}n(n^2-1) - \frac{1}{6}n(n+1)(2n+1) \\
&= \frac{1}{12}n(n^2-1) + \sum_{i=1}^n x_i y_i - \frac{1}{2} \sum_{i=1}^n x_i^2 - \frac{1}{2} \sum_{i=1}^n y_i^2 \\
&= \frac{1}{12}n(n^2-1) - \frac{1}{2} \sum_{i=1}^n x_i^2 - 2 \sum_{i=1}^n x_i y_i + \sum_{i=1}^n y_i^2 \\
&= \frac{1}{12}n(n^2-1) - \frac{1}{2} \sum_{i=1}^n (x_i - y_i)^2
\end{aligned}$$

So, since $r = \frac{12S_{xy}}{n(n^2-1)}$,

$$\begin{aligned}
r &= \frac{12 \left[\frac{1}{12}n(n^2-1) - \frac{1}{2} \sum_{i=1}^n (x_i - y_i)^2 \right]}{n(n^2-1)} \\
&= \frac{n(n^2-1) - 6 \sum_{i=1}^n (x_i - y_i)^2}{n(n^2-1)} \\
&= 1 - \frac{6 \sum_{i=1}^n (x_i - y_i)^2}{n(n^2-1)} = 1 - \frac{6 \sum d^2}{n(n^2-1)}
\end{aligned}$$

Binomial distribution formula

This formula comes from the fact that there are $\binom{n}{r}$ ways to have r out of n trials be “successful”, and the probability of each way is clearly $p^r(1-p)^{n-r}$ since clearly we need “success” r times and “failure” $n-r$ times. So the final probability becomes $\binom{n}{r}p^r(1-p)^{n-r}$.

Poisson distribution formula

Lemma: $\lim_{x \rightarrow \infty} \left(1 + \frac{a}{x}\right)^{bx} = e^{ab}$

Proof of lemma: $\ln(\lim_{x \rightarrow \infty} (1 + \frac{a}{x})^{bx}) = \lim_{x \rightarrow \infty} (\ln((1 + \frac{a}{x})^{bx}))$ (Since $\ln$ is continuous so intuitively it should therefore commute with limits. This is a standard result in analysis but since it's intuitively true I won't bother proving it)

This limit is equal to $\lim_{x \rightarrow \infty} (bx \ln(1 + \frac{a}{x})) = b \lim_{\frac{1}{x} \rightarrow 0} (\frac{\ln(1 + \frac{a}{x})}{\frac{1}{x}}) = ab \lim_{\frac{1}{x} \rightarrow 0} (\frac{\ln(1 + \frac{a}{x})}{\frac{a}{x}}) = ab \lim_{\frac{1}{x} \rightarrow 0} (\frac{\ln(1 + \frac{a}{x}) - \ln(1)}{\frac{a}{x} - 0})$ since

$\ln(1)$ is 0. But this is just ab times the derivative of $\ln(x)$ when x is 1, which is just ab . Since ab is the natural logarithm of the limit, the limit is e^{ab} as required. Now we can derive the poisson formula, by considering a scenario where the expected number of times for an event to occur in a time interval stays constant but the amount of chances for the event to happen gets larger and the probability of it happening in each time interval gets smaller, and taking a limit. This is what we get. We assume that h is n divided by an integer so that the binomial distribution in question is actually defined.

Imagine making a binomial distribution with arbitrarily short stretches of time but the same mean time period

$$X \sim \lim_{h \rightarrow 0} B(\frac{n}{h}, ph) \quad [np = \lambda]$$

$$P(X = k) = \lim_{h \rightarrow 0} \binom{\frac{n}{h}}{k} (ph)^k (1-ph)^{\frac{n}{h}-k}$$

$$= \lim_{h \rightarrow 0} \frac{(\frac{n}{h})!}{k! (\frac{n}{h}-k)!} p^k h^k (1-ph)^{\frac{n}{h}-k}$$

$$= \lim_{h \rightarrow 0} \frac{(\frac{n}{h})!}{k!} p^k h^k (1-ph)^{\frac{n}{h}-k}$$

note that $\frac{(\frac{n}{h})!}{(\frac{n}{h}-k)!}$ approaches $(\frac{n}{h})^k$

because clearly

$$\frac{(\frac{n}{h}-k)! (\frac{n}{h}-k)^k}{(\frac{n}{h}-k)!} < \frac{(\frac{n}{h})!}{(\frac{n}{h}-k)!} < \frac{(\frac{n}{h}-k)! (\frac{n}{h})^k}{(\frac{n}{h}-k)!}$$

so

$$\frac{(\frac{n}{h})!}{(\frac{n}{h}-k)!} \text{ is bounded between } (\frac{n}{h}-k)^k \text{ and } (\frac{n}{h})^k \text{ which}$$

The ratio $\frac{(\frac{n}{h}-k)^k}{(\frac{n}{h})^k}$ approaches 1 as $\frac{n}{h}$ grows so the ratio $\frac{(\frac{n}{h})!}{(\frac{n}{h}-k)!} : (\frac{n}{h})^k$ approaches 1 so the error in approximating $\frac{(\frac{n}{h})!}{(\frac{n}{h}-k)!}$ as $(\frac{n}{h})^k$ gets arbitrarily small as $\frac{n}{h}$ grows so it is valid.

Continuing simplifying $P(X = k)$,

$$= \lim_{h \rightarrow 0} \frac{(\frac{n}{h})^k}{k!} (1-ph)^{\frac{n}{h}-k} = \lim_{h \rightarrow 0} \frac{\lambda^k}{k!} (1-ph)^{\frac{n}{h}-k}$$

as $h \rightarrow 0$ this $\rightarrow 1$ so it vanishes

this becomes $e^{-\lambda}$ as $h \rightarrow 0$ which completes our proof

Image: Shows my derivation on paper

So our final formula is $\frac{e^{-\lambda} \lambda^k}{k!}$ As required.

In fact, since the poisson probabilities must sum to 1, multiplying by e^λ gives another proof for the exponential function taylor series.

Continuous distribution expected value intuition

The reason you can think of the expected value of a continuous random variable X with probability density function $f(x)$ as $\int_{-\infty}^{\infty} xf(x)dx$ is because indeed the intuition of expectation as “the sum over all possible values of the value times how likely you are to get that value” is essentially being taken to the limit as the values get closer together and the probabilities get smaller for each one (ie, $dx \rightarrow 0$). It is for the same reason that $\int_{-\infty}^{\infty} g(x)f(x)dx$. There isn't really a mathematical statement we need to prove here, it's just a definition that we are providing an intuition for.

Geometric/Negative binomial distribution mean and variance

Mean of geometric should be intuitive. Although I will give a formal proof, think: If something has a $1/3$ chance of happening, it should take on average 3 tries for it to happen!

Mean of geometric variable x with probability p ($0 < p < 1$) = $\sum_{n=0}^{\infty} nP(x = n) = \sum_{n=0}^{\infty} np(1-p)^{n-1}$ (This is intuitive, it has to happen n times and happen 1 time) = $p \sum_{n=0}^{\infty} n(1-p)^{n-1}$.

This sum is equivalent to adding the rows of the following:

$$\begin{aligned}
 &1 \\
 &+ (1-p) + (1-p) \\
 &+ (1-p)^2 + (1-p)^2 + (1-p)^2 \\
 &+ (1-p)^3 + (1-p)^3 + (1-p)^3 + (1-p)^3 \\
 &\vdots
 \end{aligned}$$

Image: Shows a sum rewritten so that all the

terms with a particular exponent are on the same row

We now consider instead adding the columns rather than the rows.

Technical note: Note that this does converge absolutely, since all terms are non-negative and the sum $\sum_{n=0}^{\infty} n(1-p)^{n-1}$ converges by the ratio test. For details on what this means, see the appendix with technical justification of the generalized binomial theorem in the A level maths supplementary material document. The reason this matters is that this means we are allowed to change the order in which we sum terms.

Now we add the columns of the image above rather than the rows, giving us that the mean of a geometric random variable with probability p is now equal to $p \sum_{n=0}^{\infty} \sum_{m=0}^{\infty} (1-p)^n (1-p)^m = p \sum_{n=0}^{\infty} \frac{(1-p)^n}{1-(1-p)}$ (geometric series) = $p \sum_{n=0}^{\infty} \frac{(1-p)^n}{p} = \sum_{n=0}^{\infty} (1-p)^n = \frac{1}{1-(1-p)}$ (geometric series) = $\frac{1}{p}$ as required.

Mean of negative binomial follows from additive property of means, since a negative binomial distribution is the sum of geometric distributions. Variance of negative binomial will follow from additive property of variances when we prove the variance of a geometric distribution, which we will do below:

We know $E[X] = \frac{1}{p}$. Then the variance is:

$$E[X^2] - (E[X])^2 = E[X(X-1)] + E[X] - (E[X])^2 = \boxed{E[X(X-1)]} + \frac{1}{p} - \frac{1}{p^2}$$

Split $E[X^2]$ into $E[X(X-1)] + E[X]$, which is easier to determine. To determine $\boxed{E[X(X-1)]}$ we have to determine the value of the following series for $p \in (0, 1)$:

$$\sum_{k=1}^{\infty} k(k-1)p(1-p)^{k-1}$$

Now we substitute $q=1-p$

$$\begin{aligned}
\sum_{k=1}^{\infty} k(k-1)p(1-p)^{k-1} &= p \sum_{k=1}^{\infty} k(k-1)(1-p)^{k-1} \\
&= p \sum_{k=1}^{\infty} (k-1)kq^{k-1} \\
&= p \frac{d}{dq} \left(\sum_{k=1}^{\infty} (k-1)q^k \right) \\
&= p \frac{d}{dq} \left(q^2 \sum_{k=1}^{\infty} (k-1)q^{k-2} \right) \\
&= p \frac{d}{dq} \left(q^2 \sum_{k=2}^{\infty} (k-1)q^{k-2} \right) \\
&= p \frac{d}{dq} \left(q^2 \frac{d}{dq} \left(\sum_{k=2}^{\infty} q^{k-1} \right) \right) \\
&= p \frac{d}{dq} \left(q^2 \frac{d}{dq} \left(\sum_{k=1}^{\infty} q^k \right) \right)
\end{aligned}$$

Justification for differentiating infinite power series inside their radius of convergence (Which this is, the ratio between consecutive terms approaches $1-p$ which is between 0 and 1) is given in the A level maths supplementary material document under the generalized binomial theorem formal justification appendix.

$$\begin{aligned}
&= p \frac{d}{dq} \left(q^2 \frac{d}{dq} \left(\frac{1}{1-q} - 1 \right) \right) \\
&= p \frac{d}{dq} \left(\frac{q^2}{(1-q)^2} \right) \\
&= p \left(\frac{-2q}{(q-1)^3} \right) \quad \text{Backsub. } q = (1-p) \\
&= p \left(\frac{-2(1-p)}{((1-p)-1)^3} \right) = p \left(\frac{-2+2p}{-p^3} \right) \\
&= \frac{-2+2p}{-p^2} = \frac{2(p-1)}{-p^2} = \frac{2(1-p)}{p^2}.
\end{aligned}$$

Now putting the result back into the equation for $Var[X]$ gives us:

$$\begin{aligned}
Var[X] &= E[X(X-1)] + E[X] - (E[X])^2 = \frac{2(1-p)}{p^2} + \frac{1}{p} - \frac{1}{p^2} = \frac{2-2p+p-1}{p^2} \\
&= \frac{1-p}{p^2}.
\end{aligned}$$

Generating functions multiply when variables add

I will demonstrate why this works with an example. Suppose we have the following random variables x and y :

K	0	1	2	3
$P(x=K)$	0.1	0.2	0.3	0.4

K	0	1	2
$P(y=K)$	0.3	0.5	0.2

Then the generating function of x is $0.1 + 0.2t + 0.3t^2 + 0.4t^3$ and the generating function of y is $0.3 + 0.5t + 0.2t^2$. Now consider how we would find the following:

1. The t^4 coefficient of the product of the generating functions
2. $P(x+y=4)$

For the first one, when we expand it, we will have a t^4 term from the $0.3t^2$ from the first term times the $0.2t^2$ from the second term, and also the $0.4t^3$ term from the first term times the $0.5t$ from the second term. The products that will give a t^4 term are exactly those whose exponents add up to 4.

Now, to find $P(x+y=4)$, consider all the possible cases: Either $x=2$ and $y=2$ or $x=3$ and $y=1$. So, since x and y are assumed to be independent, we have that $P(x+y=4)=P(x=2)P(y=2)+P(x=3)P(y=1)$. Crucially, multiplying the probabilities that the variable equals k is the same as multiplying the coefficients of t^k by the definition of generating functions, and we do it on exactly those where $x+y$, which corresponds to the sum of the exponents, equals 4.

Hopefully this is convincing enough that the product identity holds.

Linear coding property of generating functions

We want to prove that if a random variable X has probability generating function $g(t)$ then $aX+b$ has probability generating function $t^b g(t^a)$. The reason this is true is because when we have aX all the possible values for the variables are multiplied by a and these correspond to the exponents in the generating function which therefore must be multiplied by a . Then by adding b we add b to all the possible values for the variables which corresponds to adding b to the exponents in the generating function, thus it is exactly like multiplying each term by t^b .

Negative binomial generating function

We will first prove that $\sum_{x=r}^{\infty} \binom{x-1}{r-1} q^{x-r} = (1-q)^{-r}$ by induction on r , noting that $q=1-p$ in this context. Note that the differentiation of the power series we will do is allowed since in this context $|q|<1$ so we are inside the radius of convergence by the ratio test and the fact that the result of a ratio test is the same under differentiation (See the appendix on the binomial theorem in the A level maths document to see what I'm talking about).

$$\sum_{x=r}^{\infty} \binom{x-1}{r-1} q^{x-r} = (1-q)^{-r} \quad (*)$$

If $r=1$ this just becomes

$$\sum_{x=1}^{\infty} \binom{x-1}{0} q^{x-1} \text{ since } \binom{x-1}{0} = 1 \text{ so}$$

this is a geometric series which equals $(1-p)^{-1}$
 this verifies the base case

Now do induction

Differentiate $(*)$ noting that the $x=r$ term is constant so its derivative is 0.

$$r(1-q)^{-r-1} \frac{d}{dq} \left(\sum_{x=r}^{\infty} \frac{(x-1)!}{(x-r)!(r-1)!} q^{x-r} \right) = \sum_{x=r}^{\infty} \frac{(x-r)(x-1)!}{(x-r)!(r-1)!} q^{x-r-1}$$

$$= r \left(\sum_{x=r+1}^{\infty} \frac{(x-1)!}{(x-r)!(r-1)!} q^{x-r-1} \right)$$

$$\downarrow \binom{x-1}{(r-1)+1}$$

cancelling rs gives $(1-q)^{-r-1} = \sum_{x=r+1}^{\infty} \binom{x-1}{(r+1)-1} q^{x-(r+1)}$

Induction done

Image: Shows my proof by induction.

Now we have the lemma, so we will find the generating function of the negative binomial, using $q=1-p$. Note that here the sums are over x .

$$P(X = x) = \binom{x-1}{r-1} p^r q^{x-r}, \quad x = r, r+1, r+2, \dots$$

$$\begin{aligned}
 G_X(t) &= \sum \binom{x-1}{r-1} p^r q^{x-r} t^x \\
 &= \sum \binom{x-1}{r-1} p^r q^{x-r} t^x \\
 &= \sum \binom{x-1}{r-1} (pt)^r (qt)^{x-r} \\
 &= (pt)^r \sum \binom{x-1}{r-1} (qt)^{x-r} \\
 &= (pt)^r (1-qt)^{-r} \\
 &= \left(\frac{pt}{1-qt} \right)^r = \left(\frac{pt}{1-(1-p)t} \right)^r \text{ as required.}
 \end{aligned}$$

Normal + Normal = Normal

In my proof of the area under a normal distribution curve earlier, I briefly discussed the idea that the normal distribution is special since it is completely rotationally symmetrical if you plot in a higher dimensional space the probability density function of different standard normals. Suppose X and Y are $N(0,1)$ random variables that are independent, then it suffices to show that $aX+bY=N(0, a^2 + b^2)$ since if X and Y were shifted by constants the sum would just shift accordingly and it would all be fine. Here is the joint probability density function of X and Y visualised:




Image: another 2d normal plotted

on a 3d graph

As you can see, there is rotational symmetry. We actually just need to show that $aX+bY$ is normal then we will know from mean and variance additive properties the desired result. Suppose, for example, we want to find the probability density function of $P(3x+2y)$, then the probability density that $P(3x+2y=k)$, intuitively, corresponds to the area under the slice of the above diagram corresponding to $3x+2y=k$. In the end, we will have to rescale the probability density function of $3x+2y$ so that the area under that curve is 1, but it's quite obvious we have proportionality, which is what we actually need. Below is the slices I mean for some values of k .




Image: The same 2d normal with parallel

vertical planes on the same graph

Visually you can see from the rotational symmetry of the image that the function of the area under the red in each of these planes, which as discussed is what we need, is a normal distribution, completing what is not really a proof but rather a visual argument.

Unbiased variance estimator, vector approach




Image: Diagram

to show pythagoras vector approach for variance estimator

Residual sum of squares formula

The [least-squares regression line](#) is given by

$$y = ax + b,$$

where $b = \bar{y} - a\bar{x}$ and $a = \frac{S_{xy}}{S_{xx}}$, where $S_{xy} = \sum_{i=1}^n (\bar{x} - x_i)(\bar{y} - y_i)$ and $S_{xx} = \sum_{i=1}^n (\bar{x} - x_i)^2$.

Therefore,

$$\begin{aligned} \text{RSS} &= \sum_{i=1}^n (y_i - f(x_i))^2 = \sum_{i=1}^n (y_i - (ax_i + b))^2 = \sum_{i=1}^n (y_i - ax_i - \bar{y} + a\bar{x})^2 \\ &= \sum_{i=1}^n (a(\bar{x} - x_i) - (\bar{y} - y_i))^2 = a^2 S_{xx} - 2aS_{xy} + S_{yy} = S_{yy} - aS_{xy} = S_{yy} \left(1 - \frac{S_{xy}^2}{S_{xx}S_{yy}}\right) \end{aligned}$$

where $S_{yy} = \sum_{i=1}^n (\bar{y} - y_i)^2$.

The [Pearson product-moment correlation](#) is given by $r = \frac{S_{xy}}{\sqrt{S_{xx}S_{yy}}}$; therefore, $\text{RSS} = S_{yy}(1 - r^2)$.

Note: the fact that $a^2 S_{xx} = aS_{xy}$ is because of the definition of a .

Characteristic functions

Here we will define characteristic functions and prove some basic properties of them (although these basic properties have not such basic proofs). These will play a “central” role in our proof. If I have a probability distribution x , then I define the characteristic function of x as $\phi_x(t) = E(e^{itx})$. This is related to the fourier transform of the ditstribution, that is the cf divided by 2π . What we will prove here is that if you know the characteristic function of a distribution you can reverse engineer what the distribution was. To fully understand the proof, you will need to know the results in the “Some technical results” document, which is a separate document because the results are there are not only relevant to statistics, so it would be weird to include them in here. The idea will be to show that from the characteristic function you can reverse engineer the integral from a to b of our probability distribution for any a and b , which determines the entire distribution.

If f is our probability distribution, then

$\phi_f(t) = E(e^{itx}) = \int_{-\infty}^{\infty} f(y)e^{ity} dy$ is the characteristic function

And define, for arbitrary a and b with $a < b$,

$$I_{\varepsilon} := \frac{1}{2\pi} \int_{-\infty}^{\infty} \phi_f(t) e^{-\varepsilon t^2} \frac{e^{-iat} - e^{-ibt}}{it} dt$$

This is equal to

$$\frac{1}{2\pi} \int_{-\infty}^{\infty} \left[\int_{-\infty}^{\infty} f(y)e^{ity} dy \right] e^{-\varepsilon t^2} \frac{e^{-iat} - e^{-ibt}}{it} dt$$

Now we will show that

$$\frac{1}{2\pi} \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} \left| f(y)e^{ity} e^{-\varepsilon t^2} \frac{e^{-iat} - e^{-ibt}}{it} \right| dy dt$$

Is finite so that we can swap the integrals around. Note that for each fixed t ,

$$\int_{-\infty}^{\infty} \left| f(y)e^{ity} e^{-\varepsilon t^2} \frac{e^{-iat} - e^{-ibt}}{it} \right| dy = \int_{-\infty}^{\infty} |f(y)e^{ity}| |e^{-\varepsilon t^2}| \left| \frac{e^{-iat} - e^{-ibt}}{it} \right| dy =$$

$$|e^{-\varepsilon t^2}| \left| \frac{e^{-iat} - e^{-ibt}}{it} \right| \int_{-\infty}^{\infty} |f(y)e^{ity}| dy \leq |e^{-\varepsilon t^2}| \left| \frac{e^{-iat} - e^{-ibt}}{it} \right| \left| \int_{-\infty}^{\infty} f(y)e^{ity} dy \right|$$

Since this third term equals $E(e^{ity})$ which is the expectation of things not outside the complex unit circle, its absolute value must be no greater than 1. Therefore, $\int_{-\infty}^{\infty} |f(y)e^{ity} e^{-\varepsilon t^2} \frac{e^{-iat} - e^{-ibt}}{it}| dy \leq |e^{-\varepsilon t^2}| \left| \frac{e^{-iat} - e^{-ibt}}{it} \right|$, so it remains to show that $\int_{-\infty}^{\infty} |e^{-\varepsilon t^2}| \left| \frac{e^{-iat} - e^{-ibt}}{it} \right| dt$ is finite (as the factor of 2π does not change finiteness), which simplifies our problem massively.

To do this, I will put some bounds on the term $\left| \frac{e^{-iat} - e^{-ibt}}{it} \right|$. This term is equal to $\frac{|e^{-iat} - e^{-ibt}|}{|t|}$.

Now I will use the inequality $|e^{ix} - e^{iy}| < |x - y|$ if x and y are real numbers not equal to each other. The reason this inequality is true can be shown geometrically:




Image: A unit circle on the complex plane with 2 points marked to demonstrate the “shortest distance” principle

Here, the green and blue points represent values of e^{ix} for certain values of x . The distance that they are apart along the unit circle is exactly the difference between these values of x , but the distance

between these points is less than this difference, because you could get further by walking in a straight line, as that is more efficient. That is why this inequality holds.

So, if $|t| < 1$, our expression is less than $\frac{|at-bt|}{|t|}$ which is just $|a-b|$. If $|t| \geq 1$, then the numerator is the absolute value of the difference between two points on the unit circle, which cannot be more than the diameter of the unit circle which is 2, so it is at most $\frac{2}{|t|}$. So, the integral $\int_{-\infty}^{\infty} |e^{-\varepsilon t^2}| \left| \frac{e^{-iat} - e^{-ibt}}{it} \right| dt$ is at most this:

$$\int_{|t| \leq 1} |b-a| e^{-\varepsilon t^2} dt + \int_{|t| > 1} \frac{2e^{-\varepsilon t^2}}{|t|} dt$$

The first term is at most $2|b-a|$. The second term will just get larger if you remove the denominator, and then it turns into an integral which we have already proven has a finite value, which we could give in terms of the square root of pi.

Now we get this after swapping the integrals and combining some exponents:

$$I_{\varepsilon} = \int_{-\infty}^{\infty} f(y) K_{\varepsilon}(y) dy \text{ where } K_{\varepsilon}(y) = \frac{1}{2\pi} \int_{-\infty}^{\infty} e^{-\varepsilon t^2} \frac{e^{it(y-a)} - e^{it(y-b)}}{it} dt$$

Definition: $\text{erf}(x) := \frac{2}{\sqrt{\pi}} \int_0^x e^{-t^2} dt$. This function goes to 1 as x goes to infinity and -1 as x goes to negative infinity. It is essentially a cdf of the normal rescaled to go from -1 to 1.

I will now define another new function as follows:

$$H_{\varepsilon}(u) = \frac{1}{2\pi} \int_{-\infty}^{\infty} e^{-\varepsilon t^2} \frac{e^{itu} - 1}{it} dt$$

Note that $H_{\varepsilon}(u) = \frac{1}{2\pi} \int_{-\infty}^{\infty} e^{-\varepsilon t^2} \int_0^u e^{its} ds dt$.

We want to show that $\int_{-\infty}^{\infty} \int_0^u |e^{-\varepsilon t^2} e^{its}| ds dt$ is finite so that we can swap the integrals around.

$$\int_{-\infty}^{\infty} \int_0^u |e^{-\varepsilon t^2} e^{its}| ds dt = \int_{-\infty}^{\infty} \int_0^u |e^{-\varepsilon t^2}| |e^{its}| ds dt = \int_{-\infty}^{\infty} |e^{-\varepsilon t^2}| \int_0^u |e^{its}| ds dt = \int_{-\infty}^{\infty} e^{-\varepsilon t^2} \int_0^u 1 ds dt = \int_{-\infty}^{\infty} u e^{-\varepsilon t^2} dt.$$

If we use the substitution $v = t\sqrt{\varepsilon}$ it will follow that this integral is $u \sqrt{\frac{\pi}{\varepsilon}}$ which is finite. Therefore we get that $H_{\varepsilon}(u)$ is equal to this:

$$\int_0^u \left(\int_{-\infty}^{\infty} e^{-\varepsilon t^2} e^{ist} dt \right) ds$$

We now complete the square as follows:

$$-\varepsilon t^2 + ist = -\varepsilon \left(t - \frac{is}{2\varepsilon} \right)^2 - \frac{s^2}{4\varepsilon}$$

To get

$$\int_{-\infty}^{\infty} e^{-\varepsilon t^2} e^{ist} dt = e^{-\frac{s^2}{4\varepsilon}} \int_{-\infty}^{\infty} e^{-\varepsilon \left(t - \frac{is}{2\varepsilon}\right)^2} dt = \sqrt{\frac{\pi}{\varepsilon}} e^{-\frac{s^2}{4\varepsilon}}$$

The reason we can shift by an imaginary constant step and still get the same result is because our function has a global single valued antiderivative defined by the Taylor series which converges everywhere, so you can integrate along contours like these and get the same result



Image: A contour that is the reals shifted by an imaginary constant for a large interval

As we make the red part longer, we know from the shape of the bell curve that the area of the tails will go to 0, so the limit of these integrals which are all the same as the integral along the real line will be the same as the integral along the real line and be equal to the integral with the imaginary shift.

Therefore, $H_\varepsilon(u)$ is equal to $\frac{1}{2\pi} \sqrt{\frac{\pi}{\varepsilon}} \int_0^u e^{-\frac{s^2}{4\varepsilon}} ds$, which can be shown by the substitution $v = \frac{s}{2\sqrt{\varepsilon}}$ to be equal to $\frac{1}{2} \operatorname{erf}\left(\frac{u}{2\sqrt{\varepsilon}}\right)$.

Now we have the following identity from the definitions:

$$K_\varepsilon(y) = H_\varepsilon(y - a) - H_\varepsilon(y - b) = \frac{1}{2} \left(\operatorname{erf}\left(\frac{y - a}{2\sqrt{\varepsilon}}\right) - \operatorname{erf}\left(\frac{y - b}{2\sqrt{\varepsilon}}\right) \right)$$

Now consider what the K function approaches as ε approaches 0: If $y < a$, then y also is $< b$ since we defined $a < b$. Therefore, both terms will go to -1 as the inputs to the erf functions will go to -infinity, so the whole thing will go to 0. Same if $y > b$. But, if $a < y < b$, then the first term's input will go to infinity so the first term will go to 1, and the second terms input will go to -infinity so the whole thing will go to $\frac{1}{2}(1 - (-1))$ which is 1. Therefore, as ε goes to 0, K approaches the function that is 1 when y is between a and b and 0 otherwise, and it approaches $\frac{1}{2}$ at exactly a and b , but this doesn't really matter. What does matter is that K is always between -1 and 1.

Recall, we had this:

$$I_\varepsilon = \int_{-\infty}^{\infty} f(y) K_\varepsilon(y) dy \text{ where } K_\varepsilon(y) = \frac{1}{2\pi} \int_{-\infty}^{\infty} e^{-\varepsilon t^2} \frac{e^{it(y-a)} - e^{it(y-b)}}{it} dt$$

We know that since f is always positive and the integral of f over the reals is 1 since f is a probability distribution, and K is between 0 and 1, define $f_n(y) := f(y) K_{\frac{1}{n}}(y)$, then $f_n(y)$ is bounded above in absolute value by $f(y)$ which has a finite integral, so dominated convergence applies. Since $f_n(y)$ converges pointwise to f when $a < y < b$ and 0 otherwise (and what happens at exactly a and b does not affect the value of the integrals), we know by dominated convergence on f_n that as ε (or $1/n$, same

thing) goes to 0, I_ε goes to $\int_a^b f(x)dx$. But I_ε was defined only in terms of the characteristic function. So done.

We will now prove that if the characteristic function of a sequence of probability distributions converges, then the distributions converge to the probability distribution with the limit characteristic function. To do this, we will define yet another function. Let $L_\varepsilon(t) := \frac{e^{-\varepsilon t^2}}{2\pi} \frac{e^{-iat} - e^{-ibt}}{it}$, then $I_\varepsilon = \int_{-\infty}^{\infty} \phi(t) L_\varepsilon(t) dt$. L_ε is integrable with the same proof as earlier and bounded by $\frac{b-a}{2\pi}$, so if our sequence of characteristic functions $\phi_n(t)$ converges pointwise to $\phi(t)$ then $\phi_n(t) L_\varepsilon(t)$ converges to $\phi(t) L_\varepsilon(t)$ as n goes to infinity. Also, $|\phi_n(t)| \leq 1$ from earlier, so $|\phi_n(t) L_\varepsilon(t)| \leq |L_\varepsilon(t)|$. Therefore all the hypotheses for the dominated convergence theorem apply, so the limit of $\int_{-\infty}^{\infty} \phi_n(t) L_\varepsilon(t) dt$ is indeed I_ε . Therefore, letting ε approach 0, we have that the probability our distribution lands between a and b approaches that probability for a pdf with cf ϕ if the cf's converge to ϕ . So done.

Note: This means the distribution converges if the cf converges, in the sense that the weight of the distribution on any interval converges, but the probability density function need not converge pointwise. Convergence in distribution actually means the cumulative distribution function converges pointwise at all points where it is continuous, which this satisfies since we showed that all integrals of the pdf converge so the cdf converges.

How the properties above apply to discrete distributions

Note on this: A valid probability distribution need not have a well defined pdf everywhere, but we have been assuming this. It may have atoms, which are values a such that $P(x=a) > 0$. Say $P(x=3)=0.2$ but otherwise x has a normal probability density function g . Then we say that $f(x) = g(x) + 0.2\delta(x-3)$ where $\delta(x)$ is not a function that is defined in the traditional sense, but rather a function that when you integrate it returns 0 if your integration interval does not contain 0 and 1 otherwise. Ie, the following hold:

$$\int_{\mathbb{R}} \delta(x) h(x) dx = h(0)$$

where this notation means integrating over all real numbers, ie minus infinity to infinity

$$\int_a^b \delta(x) dx = \begin{cases} 1, & 0 \in [a, b], \\ 0, & \text{otherwise.} \end{cases}$$

This is called the dirac delta function. You can think of it as a spike, or as the limit of normals with mean 0 and tinier variances approaching 0. The integration theory we have used to get to this proof is still logically sound even with dirac deltas allowed. (See technical detail 2) This also allows us to work with continuous random variables, discrete random variables, or combinations of those in one unified framework. We can avoid issues by making sure we pick our a 's and b 's so that at the ends of an interval, or (if we wish to generalize this to multiple variables) at the corners of our rectangle, the cdf is continuous, so we can integrate inside the rectangle and recover the cdf.

Technical detail: This kind of "ensuring continuity" thing which we do a lot here is possible because the cdf is increasing so it can't be discontinuous on an entire interval or it wouldn't be bounded on

that interval. This is because of the thing in the technical results document that says we cannot add up uncountably many finite numbers and get a finite number. Since there is no interval where our thing is discontinuous, if we end up at a discontinuity we can just move a really really tiny as tiny as we want amount then be fine.

Technical detail 2: The monotone convergence theorem still applies to these “mixed” integrals because the Dirac-delta and continuous parts converge separately, so the usual proof carries over unchanged. Everything built on this theorem therefore remains valid.

- Dirac-delta part: In each step of our increasing sequence of functions, the total mass of the delta spikes is non-decreasing and bounded by 1, hence convergent.
- Continuous part: The continuous (normal-type) component converges in the ordinary way, giving the correct limit.
- Putting them together: It remains to justify that the total contribution from the delta spikes in the limit equals the limit of their partial sums. If there are only finitely many spikes this is immediate. If there are infinitely many, we choose any finite subset whose combined mass is within any $\varepsilon/2$ of the full limit mass. The contributions from these finitely many spikes in the sequence eventually come within $\varepsilon/2$ of their limiting total, and the remaining spikes contribute at most $\varepsilon/2$, so we can eventually get arbitrarily close to the limit we want, ie within ε for any ε we choose.

Therefore the sum of the Dirac-delta parts in the limit agrees with the limit of the sums, so the monotone convergence theorem holds exactly as in the purely continuous case.

Central limit theorem (Normal approximation is a corollary of this)

We will prove this for distributions that have a probability density function that may include dirac deltas and have finite variance - this is not all distributions, but it applies to:

- Binomial
- Poisson
- Normal
- Geometric
- Negative binomial
- Anything that comes up in the real world
- Everything we use for the chi squared proof that relies on CLT

This is exactly the kind of distributions we were implicitly assuming we had when we worked with characteristic functions. The cantor distribution is an example of something that does not work due to not having a pdf, and the cauchy distribution is an example of something that does not work due to not having finite variance – In fact for the cauchy distribution the CLT is not even true.

Setup: Let $X_1, X_2, \dots$ be identically distributed and independent where each has mean 0 and variance σ^2 which is finite.

Define the normalized sum

$$S_n = \frac{X_1 + \dots + X_n}{\sigma\sqrt{n}},$$

Note that if our variables do not have mean 0 we can shift them so that they do and still apply this argument. S_n has mean 0 and variance 1 so it is standardized, and the goal is to prove that as n gets large, S_n is well approximated by a standard normal, as this is what the central limit theorem says. The characteristic function of S_n is $E\left(e^{it\left(\frac{X_1}{\sigma\sqrt{n}}\right) + it\left(\frac{X_2}{\sigma\sqrt{n}}\right) + \dots + it\left(\frac{X_n}{\sigma\sqrt{n}}\right)}\right) = E\left(e^{it\left(\frac{X_1}{\sigma\sqrt{n}}\right)}\right) E\left(e^{it\left(\frac{X_2}{\sigma\sqrt{n}}\right)}\right) \dots E\left(e^{it\left(\frac{X_n}{\sigma\sqrt{n}}\right)}\right)$ by properties of exponents and the fact that with independent things we can split expectation and product. Therefore, $\phi_{S_n}(t) = \left(\phi_x\left(\frac{t}{\sigma\sqrt{n}}\right)\right)^n$. Therefore, we will prove that $\phi_x(u) = 1 - \frac{\sigma^2 u^2}{2} + o(u^2)$ where $o(u^2)$ means that as u goes to 0 this gets much smaller than u^2 , ie if u is small enough this becomes smaller than an arbitrary constant times u^2 . We will then be able to put $u_n = \frac{t}{\sigma\sqrt{n}}$ to get $\phi_x(u_n) = 1 - \frac{t^2}{n} + o\left(\frac{1}{n}\right)$, where this time $o\left(\frac{1}{n}\right)$ means that if n is going to infinity this gets much smaller than $\frac{1}{n}$ by any arbitrarily large constant.

Note that here, most proofs expand e^{itx} as a taylor series, but this is not justified since after this a swap of expectation and infinite summation that is not justified normally takes place. Below is the rigorous way to do this.

Define for real y

$$R(y) := e^{iy} - 1 - iy + \frac{y^2}{2}$$

$$\int_0^1 (1-s)e^{isy} ds = \left[(1-s) \frac{e^{isy}}{iy} \right]_0^1 + \int_0^1 \frac{e^{isy}}{iy} ds = -\frac{1}{iy} + \left[\frac{e^{isy}}{(iy)^2} \right]_0^1 = -\frac{1}{iy} + \frac{e^{iy}}{(iy)^2} - \frac{1}{(iy)^2}$$

Where I have used integration by parts with $u=1-s$ and $v=\text{the exponential thingy}$

$$\text{Therefore } e^{iy} - 1 - iy = (iy)^2 \int_0^1 (1-s)e^{isy} ds$$

$$\text{So } |e^{iy} - 1 - iy| = |(iy)^2| \left| \int_0^1 (1-s)e^{isy} ds \right| \leq y^2 \int_0^1 |(1-s)e^{isy}| ds$$

$$= y^2 \int_0^1 |1-s| |e^{isy}| ds = y^2 \int_0^1 |1-s| ds = \frac{y^2}{2}.$$

Therefore,

$$|R(y)| \leq |e^{iy} - 1 - iy| + \frac{y^2}{2} \leq y^2$$

$$\text{Now notice that } (iy)^2 \int_0^1 (1-s)(e^{isy} - 1) ds = e^{iy} - 1 - iy - (iy)^2 \int_0^1 (1-s) ds$$

$$= e^{iy} - 1 - iy + y^2 \int_0^1 (1-s) ds = e^{iy} - 1 - iy + \frac{y^2}{2} = R(y)$$

So,

$$|R(y)| = \left| (iy)^2 \int_0^1 (1-s)(e^{isy} - 1) ds \right| \leq y^2 \int_0^1 (1-s) s |y| ds = \frac{|y|^3}{6}$$

Now define h for real numbers as follows:

$$h(y) = \begin{cases} R(y)/y^2, & y \neq 0, \\ 0, & y = 0. \end{cases}$$

By our two bounds on R(y), h is always bounded by 1, and since it is bounded by y/6 it goes to 0 as y goes to 0.

Now, using $E(X)=0$ and therefore that $E(x^2) = E(x^2) - E(x)^2 = \sigma^2$, we have that

$$\phi(u) = \mathbb{E}[e^{iuX}] = 1 + iu \mathbb{E}[X] - \frac{u^2}{2} \mathbb{E}[X^2] + \mathbb{E}[R(uX)] = 1 - \frac{\sigma^2 u^2}{2} + \mathbb{E}[R(uX)]$$

Now, by definition of h,

$$\frac{\phi(u) - 1 + \frac{\sigma^2 u^2}{2}}{u^2} = \mathbb{E} \left[\frac{R(uX)}{u^2} \right] = \mathbb{E} [h(uX) X^2]$$

Therefore if this goes to 0 as u goes to 0, we have $o(u^2)$ so we will be done with finding the cf of x.

We want to show:

$$\lim_{u \rightarrow 0} \frac{1}{u^2} \mathbb{E}[R(uX)] = 0$$

where

$$R(y) = e^{iy} - 1 - iy + \frac{y^2}{2}$$

and

$$\mathbb{E}[R(uX)] = \int_{\mathbb{R}} R(ux) f(x) dx$$

with f the probability density of X .

Define the **sequence of functions** (parametrized by $u \rightarrow 0$):

$$g_u(x) = \frac{R(ux)}{u^2} x^2$$

But notice from the proof:

$$\frac{1}{u^2} \mathbb{E}[R(uX)] = \mathbb{E} \left[\frac{R(uX)}{u^2} \right]$$

and we write

$$\frac{R(uX)}{u^2} = h(uX)X^2$$

where

$$h(y) = \begin{cases} \frac{R(y)}{y^2} & y \neq 0 \\ 0 & y = 0 \end{cases}$$

and for all y , $|h(y)| \leq 1$.

Parametrized by $u \rightarrow 0$ essentially means we can define a sequence of functions g_n by setting $u=1/n$, similar to what we did before. Now, since $|h(y)| \leq 1$, we have this

$$\mathbb{E}[h(uX)X^2] = \int_{\mathbb{R}} h(ux)x^2 f(x) dx$$

Where the integrand is bounded above by the function $x^2 f(x)$. This is integrable because $\text{Var}(x)$ is finite. This is an assumption of the central limit theorem. If $\text{Var}(x)$ is not finite, the central limit theorem is not always true! The cauchy distribution is an example of this. Also, since $h(y)$ goes to 0 as y goes to 0, we can do as follows:

$$h(ux) \rightarrow h(0) = 0$$

$$h(ux)x^2 \rightarrow 0.$$

So our integral converges pointwise to 0, so all the hypotheses of the dominated convergence theorem are met, so this term that we needed to go to 0 goes to 0, so done.

Ok, now to find the characteristic function of S_n , remember that it is $\left(\phi_x\left(\frac{t}{\sigma\sqrt{n}}\right)\right)^n$ so therefore since the characteristic function of x is $1 - \frac{t^2}{n} + o\left(\frac{1}{n}\right)$, the cf of S_n is $\left(1 - \frac{t^2}{n} + o\left(\frac{1}{n}\right)\right)^n$, which by the same limit we used in the poisson distribution proof, goes to $e^{-\frac{t^2}{2}}$ as n goes to infinity. Now it just remains to show that this is indeed the cf of a standard normal, and we will be done at last. The cf of a standard normal is

$$\int_{-\infty}^{\infty} f(y)e^{ity} dy$$

Where $f(y)$ is the formula for a standard normal, so we have

$$\begin{aligned} \int_{-\infty}^{\infty} \frac{e^{-\frac{y^2}{2}}}{\sqrt{2\pi}} e^{ity} dy &= \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\infty} e^{-\frac{y^2}{2}} e^{ity} dy = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\infty} e^{-\frac{1}{2}(y^2 - 2ity)} dy = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\infty} e^{-\frac{1}{2}(y-it)^2 - t^2} dy \\ &= \frac{1}{\sqrt{2\pi}} e^{-t^2} \int_{-\infty}^{\infty} e^{-\frac{1}{2}(y-it)^2} dy = e^{-t^2} \end{aligned}$$

Where we have used the same contour integration idea as in the characteristic function proof, so done.

Note: The reason we did not just use the Taylor series expansion is because we have not justified swapping expectation with an infinite sum.

Chi squared tests

The problem is defining what a degree of freedom is AND showing that the test statistic $\sum_{n=1}^k \frac{(O_n - E_n)^2}{E_n}$ is well approximated for fairly large n by a distribution that depends ONLY on this number. And in the proof we will better understand what this distribution (That you are taught is a chi squared distribution) really is.

Definition (random vector): A random vector is a vector of random variables that are possibly not independent. It may have its own probability density function of multiple variables and it always has its own cumulative distribution function.

Now the setup of a chi squared test is that we have a list of probabilities of several events determined by the null hypothesis. For example, for testing if a dice is fair, this list is $(\frac{1}{6}, \frac{1}{6}, \frac{1}{6}, \frac{1}{6}, \frac{1}{6}, \frac{1}{6})$, and in this proof we will call these $(p_1, p_2, p_3, p_4, p_5, p_6)$

Now if we roll the dice 120 times, our probability distribution becomes a lattice in 5D space with the “mean/expected” vector being $(np_1, np_2, \dots, np_k) = (20, 20, 20, 20, 20, 20)$ in this case. In general the lattice is in $(k-1)$ -D space where k is the number of cells in the table (Usually).

If we are testing something like a binomial distribution and use the data to form our null hypothesis, then we are assuming that the true mean is the same as the mean estimated by the data we got, which is effectively finding the probability distribution **given** that we have that linear constraint, and this reduces the number of dimensions that the distribution lives in by 1, provided they are independent from other existing constraints. This dimension count is exactly what we call the degrees of freedom.

Now similarly if we have a contingency table then fixing the first $r-1$ row totals adds $r-1$ constraints and fixing the last row does nothing as fixing the rest fixes the last one automatically, and similarly we have $c-1$ constraints for columns and each one decreases the dimension of 1. Since the unconstrained model has $rc-1$ degrees of freedom as there are rc cells, the constrained model has $rc-r-c+1$ degrees of freedom which is $(r-1)(c-1)$ if you expand it.

Therefore the first step of understanding what a degree of freedom is is done.

From now on I will use x_r to denote the r 'th entry of a vector x .

Now consider our contingency table or list of observed data to be a random vector X and total n and number of cells k . will consider the random vector with general r 'th entry $\sqrt{n} \left(\frac{\bar{x}_r - p_r}{\sqrt{p_r}} \right)$, where we note that the expected value of $\bar{x}_r$ is p_r (In the dice example, the number of ones divided by the number of rolls would be $\bar{x}_1$ and the expected value of this is $p_1 = \frac{1}{6}$). This begs the question of why we are specifically considering the vector given by $\sqrt{n} \left(\frac{\bar{x}_r - p_r}{\sqrt{p_r}} \right)$, which is what I call the standardized version of m , or m moved to the standardized world, where m is what I am calling the random vector of the table itself. The answer is because of what happens when you compute the square of its magnitude.

The square of the magnitude is the sum of the squares of its components. This is $\sum_{r=1}^k n \frac{(\bar{x}_r - p_r)^2}{p_r}$. Lets now multiply the numerator and denominator by n to get $\sum_{r=1}^k \frac{n^2 (\bar{x}_r - p_r)^2}{n p_r} = \sum_{r=1}^k \frac{(n \bar{x}_r - n p_r)^2}{n p_r} = \sum_{r=1}^k \frac{(O_r - E_r)^2}{E_r}$: Exactly the test statistic.

Now a chi squared distribution is defined as follows:

Take a normal distribution in k dimensions where each component of this k-dimensional random vector is $N(0,1)$ and they are all independent, then note that the probability density of this function is rotationally symmetric (This is a property of the normal distribution we discussed earlier and it comes from pythagoras – this property is why it comes up so much). A χ_k^2 distribution is defined as the distribution of the square of the absolute value of this random vector, or the sum of the squares of its components, or the sum of the squares of k $N(0,1)$ variables.

Therefore, we want to investigate $\sqrt{n} \left(\frac{\bar{x}_r - p_r}{\sqrt{p_r}} \right)$ when constrained to lattices in the plane and hope that we can show that it is standard normal so that we get the result we want. To do this, we will show that the probability of all versions of the contingency table approach what a normal distribution would predict, in a precise sense described below. It is not enough to merely show that the total-only constraint version converges to a k-1-dimensional normal and then assert that constraining a k-1 dimensional normal to s-dimensional slices gives an s-dimensional normal, as although this is true, if the actual probability distribution of tables had a small number of bad slices, the whole k-1 dimensional thing could still converge to a standard normal. This is essentially what I did in the old version of the proof, although with some bounding arguments to try to rule out this behavior which may have been enough to derive a contradiction if we had a bad slice for infinitely many n but this approach is cleaner and more robust. Therefore we need to control the probability of each table configuration, as if we do that instead of controlling just the limiting distribution which allows for probabilities of nearby tables to oscillate and average out and generally behave badly, it will be enough.

Philosophical remark: The statement “Chi squared tests work” is not very precise so we need to pick a statement that will justify that chi squared tests work. So here is that statement:

Take the vector of expected values and the nearest lattice point in the unconstrained version of the distribution (there is one at most 1 count away in each cell) and call this table m^* , and let y be $\left| \sqrt{n} \left(\frac{\bar{x} - p}{\sqrt{p}} \right) \right|$ (the “standardized” vector). Also fix a $0 < \varepsilon < \frac{1}{6}$ and a region R defined by $|x - E| < n^{\frac{1}{2} + \varepsilon}$, and let m be an arbitrary table. Here x is the counts vector-

Then we claim that

1. $p(m \in R) \rightarrow 1$ as $n \rightarrow \infty$
2. $\sup_{m \in R} \left(\left| \frac{p(m)}{p(m^*) e^{-\frac{1}{2}|y|^2}} - 1 \right| \right) \rightarrow 0$ as $n \rightarrow \infty$

Note that we are considering the ratio $\frac{p(m)}{p(m^*) e^{-\frac{1}{2}|y|^2}}$ here as the standard normal predicts that the ratio between the probability of an arbitrary table and the propbability of the peak table is about $e^{-\frac{1}{2}|y|^2}$


In english, we will show that for large enough n , the probability of any table in a region with probability approaching 1 as n gets large can be made close to the prediction of the standard normal with as narrow of a margin as we want, and also that this is true when conditioning to a lattice.

Note a subtlety here: If we merely showed that $\left| \frac{p(m)}{p(m^*)e^{-\frac{1}{2}|y|^2}} - 1 \right| \rightarrow 0$ for each m this would not be enough. I know this isn't a good example as m depends on n , but the idea is the following:

Pointwise convergence of a function (ie, for each value, the function approaches this value) and uniform convergence (ie, for any tolerance, eventually the function is within that tolerance for all values) are saying different things. I will review this definition in my level 7 differential equations notes. As an example (which is the same example as used in those notes), is this sequence of functions that converges pointwise but not uniformly:

$$f_n(x) = \begin{cases} 0 < x < \frac{1}{n}: 1 - nx \\ \text{Otherwise: } 0 \end{cases}$$

This looks like the functions in these images:



Images: Shows

examples of this for $n=1, 2, 10$.

The point is: For any point you choose, say k , after $> \frac{1}{k}$ steps, we will get to 0. So we converge pointwise to 0. It seems like we can pick $x=0$ and it will converge to 1 but we defined $f(0)$ to be 0. We do not converge uniformly to 0 as we never satisfy that all points get arbitrarily close to 0.

Note that we need uniform convergence for this proof, which is why we add in the supremum operator, because we need to show that there is a fixed n where the normal approximation is fairly good everywhere.

At this point, Claims 1 and 2 give us uniform control over the probability of every table in the central region R . Because this control holds table-by-table, we are allowed to impose a linear constraint and reason exactly as in the continuous normal case: inside R , conditioning simply restricts us to a lower-dimensional normal shape (by symmetry, as explained earlier). What is not yet guaranteed is that conditioning cannot concentrate a significant amount of probability into the tails, since Claim 1 only controls the total probability of the tails before conditioning. A third claim is therefore needed to show that the tail bound from Claim 1 still holds after conditioning, so that almost all of the conditional probability remains inside R .

3. Define R_2 by $|x - E| < An^{\frac{1}{2}+\epsilon}$ (where A is a constant that depends on the specific probabilities but not on n which we will explain how to find in the proof of the lemma) which still has

probability going to 1 by precisely the same proof as part 1 of the claim. Then the claim is that for any S intersecting R_2 , $p(m \in R \mid m \in S) \rightarrow 1$ as $n \rightarrow \infty$ where S is any lattice we constrain to: We need this because we only have control inside R by parts 1 and 2 so in principle there could be “bad lattices” where this does not hold without contradicting part 1 provided this happens with probability going to 0 as $n \rightarrow \infty$.

Since we control the probability of every single possible table as well as the behavior in every slice in the “almost certain” zone, we can say confidently once we have proven this that chi squared tests work better as n gets large as if n is large enough we will eventually have to go far into the “rejection zone” to run into issues anyways.

Note that in the unconstrained model, $|m^* - E| \leq \sqrt{k}$. This is because there exists a possible table within a distance $\sqrt{k}$ of E and the reason is because we can round everything in E down and then add 1 to each component until the total is correct, and then each component will be within 1 of E and it will be a possible table with distance at most $\sqrt{k}$ from E.

Idea of why this is all we need:

Once we prove these claims, then for a fixed circle/sphere A in the standardized world inside a lattice S where the line from its center to m^* is perpendicular to S and has length L (ie the center is near the most likely table in S) (Since we are interested in the square of the distance from the central point), we

$$\text{have } P(m \in A \mid m \in S) \approx \frac{\sum_{m \in A} P(m^*) e^{-\frac{1}{2}|y|^2 - \frac{1}{2}|L|^2}}{\sum_{m \in S} P(m^*) e^{-\frac{1}{2}|y|^2 - \frac{1}{2}|L|^2}} \text{ which is exactly } \frac{\sum_{m \in A} e^{-\frac{1}{2}|y|^2}}{\sum_{m \in S} e^{-\frac{1}{2}|y|^2}}$$

where this is off by a factor that approaches 1, and as n gets large the lattice in the standardized world gets as fine as we want since distances are being divided by $O(\sqrt{n})$ so it can be approximated by the riemann integral inside a circle/sphere of a standard normal divided by the integral over the whole thing (since it gets arbitrarily fine), and this is the probability that the distance of a normal is within that value, and we care about the distribution of the distance squared, and thus the probability that this value is less than a fixed value.

So now lets prove the claim. First we need a lemma:

Lemma 1 (stirling's approximation): For all positive integers n, $\sqrt{2\pi n} \left(\frac{n}{e}\right)^n < n! \leq \sqrt{2\pi n} \left(\frac{n}{e}\right)^n e^{\frac{1}{12n}}$

Proof of lemma 1:

$$\text{Define } d_n := \log\left(\frac{n!e^n}{n^{n+\frac{1}{2}}}\right) = \log(n!) - \left(n + \frac{1}{2}\right)\log(n) + n$$

$$\text{Then } d_n - d_{n+1} = \log(n!) - \left(n + \frac{1}{2}\right)\log(n) + n - \log((n+1)!) + \left(n + \frac{3}{2}\right)\log(n+1) - n - 1$$

$$d_n - d_{n+1} = -\log(n+1) - \left(n + \frac{1}{2}\right)\log(n) + \left(n + \frac{3}{2}\right)\log(n+1) - 1$$

$$d_n - d_{n+1} = -\left(n + \frac{1}{2}\right)\log(n) + \left(n + \frac{1}{2}\right)\log(n+1) - 1 = \left(n + \frac{1}{2}\right)\left(\log\left(\frac{n+1}{n}\right)\right) - 1$$

$$\text{Now do the substitution } t := \frac{1}{2n+1}. \text{ } n = \frac{1-t}{2t} \text{ Then } d_n - d_{n+1} = \left(\frac{1}{2} + \frac{1-t}{2t}\right)\log\left(\frac{\frac{1-t}{2t}+1}{\frac{1-t}{2t}}\right) - 1$$

$$= \frac{1}{2t} \log\left(\frac{1+t}{1-t}\right) - 1$$

Note that by the log taylor series, $\log(1+t) = t - \frac{t^2}{2} + \frac{t^3}{3} - \frac{t^4}{4} + \dots$ ($|t| < 1$ automatically by how t is defined here so it is valid), and $\log(1-t) = -\left(t + \frac{t^2}{2} + \frac{t^3}{3} + \frac{t^4}{4} + \dots\right)$.

$$\begin{aligned} \text{Therefore } \frac{1}{2t} \log\left(\frac{1+t}{1-t}\right) - 1 &= \frac{1}{2t} \log(1+t) - \frac{1}{2t} \log(1-t) - 1 \\ &= \frac{1}{2t} \left(t - \frac{t^2}{2} + \frac{t^3}{3} + \frac{t^4}{4} + \dots\right) + \frac{1}{2t} \left(t + \frac{t^2}{2} + \frac{t^3}{3} + \frac{t^4}{4} + \dots\right) - 1 \\ &= \frac{1}{2t} \left(2t + \frac{2t^3}{3} + \frac{2t^5}{5} + \frac{2t^7}{7} + \dots\right) - 1 = 1 + \frac{t^2}{3} + \frac{t^4}{5} + \frac{t^6}{7} + \dots - 1 = \frac{t^2}{3} + \frac{t^4}{5} + \frac{t^6}{7} + \dots \end{aligned}$$

t is positive so $d_n - d_{n+1} < \frac{t^2}{3} + \frac{t^4}{5} + \frac{t^6}{7} + \dots = \frac{1}{3} \frac{t^2}{1-t^2}$ (geometric series) $= \frac{1}{3((2n+1)^2-1)}$ by definition of t .

By simple algebra this is $\frac{1}{12} \left(\frac{1}{n} - \frac{1}{n+1}\right)$.

Now apply the method of differences/telescoping. $d_1 - d_n = d_1 - d_2 + d_2 - d_3 + \dots + d_{n-1} - d_n$

$$< \frac{1}{12} \left(\frac{1}{1} - \frac{1}{2}\right) + \frac{1}{12} \left(\frac{1}{2} - \frac{1}{3}\right) + \dots + \frac{1}{12} \left(\frac{1}{n-1} - \frac{1}{n}\right) = \frac{1}{12} \left(1 - \frac{1}{n}\right)$$

Note that $d_n - d_{n+1}$ is positive so d is a decreasing sequence, and for all n , d_n is bounded below by $d_1 - \frac{1}{12}$ for all n by the inequality above. A decreasing sequence that is bounded below converges and therefore our sequence d_n converges to a limit which we will call A which is also a strict lower bound for all terms d_n in our decreasing sequence.

Now let $m > n$, then $d_n - d_m < \frac{1}{12} \left(\frac{1}{n} - \frac{1}{n+1}\right) + \frac{1}{12} \left(\frac{1}{n+1} - \frac{1}{n+2}\right) + \dots + \frac{1}{12} \left(\frac{1}{m-1} - \frac{1}{m}\right) = \frac{1}{12} \left(\frac{1}{n} - \frac{1}{m}\right)$. Taking the limit as $m \rightarrow \infty$ we get that $d_n - A \leq \frac{1}{12n}$ (not strict as a sequence always strictly above a value can converge to that value). Therefore $A < d_n \leq A + \frac{1}{12n}$.

Now we will do some magic by working on a seemingly irrelevant integral where A magically pops out.

$$I_n := \int_0^{\frac{\pi}{2}} \sin^n(\theta) d\theta = [-\cos(\theta) \sin^{n-1}(\theta)]_0^{\frac{\pi}{2}} + (n-1) \int_0^{\frac{\pi}{2}} \sin^{n-2}(\theta) \cos^2(\theta) d\theta$$

Where we did integration by parts integrating $\sin$ and differentiating $\sin^{n-1}(\theta)$ and were careful to not flip any minus signs the wrong way even when we had double negatives to deal with.

$I_n = (n-1) \int_0^{\frac{\pi}{2}} \sin^{n-2}(\theta) - \sin^n(\theta) d\theta$ because the bracketed part vanishes to 0 and we have the pythagorean identity $\cos^2(\theta) = 1 - \sin^2(\theta)$. Thus $I_n = (n-1)(I_{n-2} - I_n)$ so $nI_n = (n-1)I_{n-2}$ so we have our recurrence relation $I_n = \frac{n-1}{n} I_{n-2}$ for n at least 2. Since $I_0 = \frac{\pi}{2}$ and $I_1 = 1$ by basic calculus we have the following:

$I_{2n} = \frac{1}{2} \frac{3}{4} \dots \frac{2n-1}{2n} \frac{\pi}{2}$ and $I_{2n+1} = \frac{2}{3} \frac{4}{5} \dots \frac{2n}{2n+1}$ by the recurrence relation. It is also easy to see that I_n is a decreasing sequence since the functions $\sin^n(\theta)$ are decreasing as n increases as $0 < \sin^n(\theta) < 1$

everywhere in the domain except at $\frac{\pi}{2}$. Therefore $1 \leq \frac{l_{2n}}{l_{2n+1}} \leq \frac{l_{2n-1}}{l_{2n+1}} = 1 + \frac{1}{2n}$ by the recurrence relation

$$I_{2n} = \frac{1}{2} \frac{3}{4} \cdots \frac{2n-1}{2n} \frac{\pi}{2} = \frac{1 * 2 * 3 * \dots * 2n}{2 * 2 * 4 * 4 * \dots * 2n * 2n} \frac{\pi}{2} = \frac{(2n)!}{(n! 2^n)^2} \frac{\pi}{2}$$

$$I_{2n+1} = \frac{2}{3} \frac{4}{5} \cdots \frac{2n}{2n+1} = \frac{2 * 2 * 4 * 4 * \dots * 2n * 2n}{1 * 2 * 3 * \dots * 2n * (2n+1)} = \frac{(n! 2^n)^2}{(2n+1)!}$$

Note that when n is large the ratio between $n!$ and $n^{n+\frac{1}{2}}e^{-n+A}$ converges to 1 as the log of the difference converges to 0 by the definition of A and the definition of d_n .

$$\text{Now } \frac{I_{2n}}{I_{2n+1}} = \frac{\frac{(2n)!}{(n!2^n)2^2} \pi}{\frac{(n!2^n)^2}{(2n+1)!}} = \frac{(2n)!(2n+1)!\pi}{(n!2^n)^4 2} = \pi(2n+1) \left[\frac{((2n)!)^2}{2^{4n+1}(n!)^4} \right]$$

$$\begin{aligned} \frac{((2n)!)^2}{2^{4n+1}(n!)^4} &= \frac{\left(\frac{(2n)! e^{2n}}{(2n)^{2n+\frac{1}{2}}}\right)^2}{2^{4n+1}(n!)^4} \frac{e^{-4n}}{(2n)^{-1-4n}} = \frac{e^{2d_{2n}}}{2^{4n+1} \left(\frac{n! e^n}{n^{n+\frac{1}{2}}}\right)^4} \frac{e^{4n}}{n^{4n+2}} \frac{e^{-4n}}{(2n)^{-1-4n}} = \frac{e^{2d_{2n}}}{2^{4n+1} e^{4d_n}} \frac{1}{n^{4n+2}} \frac{1}{(2n)^{-1-4n}} \\ &= \frac{e^{2d_{2n}}}{2^{4n+1} e^{4d_n}} \frac{1}{n^{4n+2}} \frac{1}{2^{-1-4n} n^{-1-4n}} = \frac{e^{2d_{2n}}}{e^{4d_n}} \frac{1}{n} \end{aligned}$$

As n gets large the ratio between this and $\frac{1}{ne^{2A}} \rightarrow 1$ and thus the limit as n goes to infinity of the ratio

$\frac{I_{2n}}{I_{2n+1}} = \pi(2n+1) \left[\frac{((2n)!)^2}{2^{4n+1}(n!)^4} \right]$ (which we know is 1) is exactly the limit of $\pi(2n+1) \frac{1}{ne^{2A}} = \frac{2\pi}{e^{2A}}$, and therefore A must be $\log(\sqrt{2\pi})$. I told you this would be magic!

Thus, we are about to get the lemma.

$$A < d_n \leq A + \frac{1}{12n}$$

$$\log(\sqrt{2\pi}) < \log\left(\frac{n! e^n}{n^{n+\frac{1}{2}}}\right) \leq \log(\sqrt{2\pi}) + \frac{1}{12n}$$

$$\sqrt{2\pi} < \frac{n! e^n}{n^{n+\frac{1}{2}}} \leq \sqrt{2\pi} e^{\frac{1}{12n}}$$

$$\sqrt{2\pi n} < \frac{n! e^n}{n^n} \leq \sqrt{2\pi n} e^{\frac{1}{12n}}$$

$$\sqrt{2\pi n} \left(\frac{n}{e}\right)^n < n! \leq \sqrt{2\pi n} \left(\frac{n}{e}\right)^n e^{\frac{1}{12n}}$$

So done (with the lemma).

Remark: $\sqrt{2\pi n} \left(\frac{n}{e}\right)^n e^{\frac{1}{12n}}$ is far closer to $n!$ in practice than the lower bound, in fact it is possible to show that $\sqrt{2\pi n} \left(\frac{n}{e}\right)^n e^{\frac{1}{12n+1}}$ is a lower bound and strict inequalities are true for both bounds but we

will not prove this as although it is not too difficult it is quite long and technical and the bound given above is sufficient for the proof we will do.

Proof of part 1 of the claim:

We will work on the unconstrained model for now since the bounds we will get on the probabilities will allow us to condition onto sub-lattices later.

Suppose $|x - E| \geq n^{\frac{1}{2}+\varepsilon}$ but for all r from 0 to k , $|O_r - E_r| < \frac{n^{\frac{1}{2}+\varepsilon}}{\sqrt{k}}$. Then we would have $|x - E| < n^{\frac{1}{2}+\varepsilon}$

by pythagoras ($|x - E| = \sqrt{\sum_{r=1}^k (O_r - E_r)^2} < \sum_{r=1}^k \frac{n^{\frac{1}{2}+\varepsilon}}{k} = \sqrt{n^{1+2\varepsilon}} = n^{\frac{1}{2}+\varepsilon}$). Hence $|x - E| \geq n^{\frac{1}{2}+\varepsilon}$

implies that there is at least one r with $|O_r - E_r| \geq \frac{n^{\frac{1}{2}+\varepsilon}}{\sqrt{k}}$. But the r 'th entry of the vector is on its own is distributed like a binomial distribution (minus the mean) so its standard deviation is proportional to $\sqrt{n}$, so the number of standard deviations away the r 'th entry of the vector would have to be from its value is proportional to n^ε which goes to infinity (slowly) as n goes to infinity.

We want to show that the probability of the r 'th component being out by that much goes to 0, as it will then follow that the probability of any of the k components being out by that much is 0 since it is a fixed number (k) of separate events each with probability going to 0. Well, it seems obvious that being out by an unbounded number of standard deviations has probability approaching 0 but we want to show this rigorously.

Therefore we will prove an easy inequality (due to markov) that says that $P(y \geq a) \leq \frac{E(y)}{a}$ where y is any distribution.

The proof is that $E(y) \geq E(y|y \geq a) \geq E(a|y \geq a) = aP(y \geq a)$ so rearranging gives the desired result.

We prove another easy inequality (due to chebyshev) that says that for any distribution x we have that $P(|x - E(x)| \geq t) \leq \frac{Var(x)}{t^2}$. The proof is simple: Just note $P(|x - E(x)|^2 \geq t^2) \leq \frac{Var(x)}{t^2}$ is equivalent to what we are claiming but it is just the above inequality applied to $|x - E(x)|^2$.

Therefore, since $Var(O_r) = np_r(1 - p_r) < n$, taking $t = \frac{n^{\frac{1}{2}+\varepsilon}}{\sqrt{k}}$. Then the inequality above gives that

$$P(|O_r - E_r| \geq t) \leq \frac{Var(x)}{t^2} < \frac{kn}{n^{1+2\varepsilon}} = kn^{-2\varepsilon} \rightarrow 0.$$

Note that from now on I use m_r, O_r interchangeably to mean the same thing.

Proof of part 2 of the claim:

We need to better understand the probability of a particular table in the unconstrained model. The idea is we can get it in terms of factorials then apply lemma 1.

Suppose the total is n and the counts are $m_1, m_2, \dots, m_k$. Then we will argue as follows:

We do n trials and the number of ways to interchange the order of these n trials is $n!$, however for each possible outcome, the order of those is irrelevant (For example, if you roll 120 dice, then we can change the order that we got the 1's in 20! Different ways without changing the overall order).

Therefore the number of possible ways to order the trials without distinguishing the same outcome is

exactly $\frac{n!}{m_1!m_2!\dots m_k!}$, and the probability of one particular ordering is $p_1^{m_1}p_2^{m_2}\dots p_k^{m_k}$. Therefore the exact probability of a particular table is $\frac{n!}{m_1!m_2!\dots m_k!}p_1^{m_1}p_2^{m_2}\dots p_k^{m_k}$.

Taking logs of lemma 1 gives $\log(m!) = \left(m + \frac{1}{2}\right)\log(m) - m + \frac{1}{2}\log(2\pi) + r_m$ with $0 < r_m < \frac{1}{12}$.

Now set $h_i := m_i - m_i^* = O_i - E_i + O(1)$. Therefore in the standardized world, $y_i = \frac{h_i}{\sqrt{np_i}} + O\left(n^{-\frac{1}{2}}\right)$.

Here when I use say $O(f(x))$ I mean there exists a constant C not depending on n or m such that for all m in R the thing we are saying is $|O(f(x))|$ is less than $Cf(x)$

$$\log(P(m)) = \log(n!) - \sum_{i=1}^k \log(m_i!) + \sum_{i=1}^k m_i \log(p_i)$$

By lemma 1,

$$\begin{aligned} \log(P(m)) &= \left(n + \frac{1}{2}\right)\log(n) - n + \frac{1}{2}\log(2\pi) + r_n - \sum_{i=1}^k \left(m_i + \frac{1}{2}\right)\log(m_i) - m_i + \frac{1}{2}\log(2\pi) + r_{m_i} + \sum_{i=1}^k m_i \log(p_i) \\ \log(P(m^*)) &= \left(n + \frac{1}{2}\right)\log(n) - n + \frac{1}{2}\log(2\pi) + r_n - \sum_{i=1}^k \left(m_i^* + \frac{1}{2}\right)\log(m_i^*) - m_i^* + \frac{1}{2}\log(2\pi) + r_{m_i^*} + \sum_{i=1}^k m_i^* \log(p_i) \end{aligned}$$

Then

$$\log(P(m)) - \log(P(m^*)) = - \sum_{i=1}^k \left(m_i + \frac{1}{2}\right)\log(m_i) - h_i + r_{m_i} + \sum_{i=1}^k h_i \log(p_i) + \sum_{i=1}^k \left(m_i^* + \frac{1}{2}\right)\log(m_i^*) + r_{m_i^*}$$

Then the remainder terms contribute $|\sum_{i=1}^k r_{m_i} - r_{m_i^*}| \leq \sum_{i=1}^k \frac{1}{12} \left| \frac{1}{m_i} - \frac{1}{m_i^*} \right|$ by the triangle inequality.

If $n > \frac{8}{p_i^3}$ for all i and we are in R, $m_i > m_i^* - n^{\frac{2}{3}} > \frac{1}{2}np_i$ which is a constant times n. Therefore, if n is large enough, $|\sum_{i=1}^k r_{m_i} - r_{m_i^*}| \leq \sum_{i=1}^k \frac{1}{12} \left| \frac{h_i}{m_i m_i^*} \right| \leq C \sum_{i=1}^k \left| \frac{h_i}{n^2} \right|$ for some fixed C. This is therefore $O\left(n^{\varepsilon - \frac{3}{2}}\right)$.

$$m_i = m_i^*(1 + u_i) \text{ with } u_i = \frac{h_i}{m_i^*} = \frac{O\left(n^{\frac{1}{2} + \varepsilon}\right)}{np_i + O(1)} = O\left(n^{\varepsilon - \frac{1}{2}}\right)$$

Write $\Delta(m) := \log(P(m)) - \log(P(m^*)) = \sum_{i=1}^k -\left(m_i + \frac{1}{2}\right)\log(m_i) + h_i + h_i \log(p_i) + \left(m_i^* + \frac{1}{2}\right)\log(m_i^*) + O\left(n^{\varepsilon - \frac{3}{2}}\right)$

Write $m_i = m_i^*(1 + u_i)$ so $u_i := \frac{h_i}{m_i^*} < \frac{n^{\frac{1}{2} + \varepsilon}}{np_i + O(1)} = O\left(n^{\varepsilon - \frac{1}{2}}\right)$.

We note that $\sum h_i = 0$ so

$$\Delta(m) = \sum_{i=1}^k -\left(m_i + \frac{1}{2}\right)\log(m_i) + h_i \log(p_i) + \left(m_i^* + \frac{1}{2}\right)\log(m_i^*) + O\left(n^{\varepsilon - \frac{3}{2}}\right)$$

Now $\left(m_i + \frac{1}{2}\right)\log(m_i) = \left(m_i^* + h_i + \frac{1}{2}\right)(\log(m_i^*) + \log(1 + u_i))$

Therefore

$$\Delta(m) = \sum_{i=1}^k -\left(m_i^* + h_i + \frac{1}{2}\right) (\log(m_i^*) + \log(1 + u_i)) + h_i \log(p_i) + \left(m_i^* + \frac{1}{2}\right) \log(m_i^*) + O(n^{\varepsilon - \frac{3}{2}})$$

$$\Delta(m) = \sum_{i=1}^k -m_i^* \log(m_i^*) - h_i \log(m_i^*) - \frac{1}{2} \log(m_i^*) - m_i^* \log(1 + u_i) - h_i \log(1 + u_i) - \frac{1}{2} \log(1 + u_i) + h_i \log(p_i) + \left(m_i^* + \frac{1}{2}\right) \log(m_i^*) + O(n^{\varepsilon - \frac{3}{2}})$$

$$\Delta(m) = \sum_{i=1}^k -h_i \log(m_i^*) - m_i^* \log(1 + u_i) - h_i \log(1 + u_i) - \frac{1}{2} \log(1 + u_i) + h_i \log(p_i) + O(n^{\varepsilon - \frac{3}{2}})$$

$$\Delta(m) = \sum_{i=1}^k -h_i \log(m_i^*) - m_i \log(1 + u_i) - \frac{1}{2} \log(1 + u_i) + h_i \log(p_i) + O(n^{\varepsilon - \frac{3}{2}})$$

$$h_i (\log(m_i^*) - \log(p_i)) = h_i \left(\log\left(\frac{m_i^*}{p_i}\right) \right) = h_i (\log(n + O(1))) = h_i (\log(n) + O\left(\frac{1}{n}\right))$$

$$\sum_{i=1}^k h_i (\log(m_i^*) - \log(p_i)) = \sum_{i=1}^k h_i \left(O\left(\frac{1}{n}\right) \right) = O\left(n^{\varepsilon - \frac{1}{2}}\right)$$

This is because $\sum h_i = 0$

Then

$$\Delta(m) = \sum_{i=1}^k -m_i^* \log(1 + u_i) - h_i \log(1 + u_i) - \frac{1}{2} \log(1 + u_i) + O(n^{\varepsilon - \frac{1}{2}})$$

$$\Delta(m) = \sum_{i=1}^k -m_i \log(1 + u_i) - \frac{1}{2} \log(1 + u_i) + O(n^{\varepsilon - \frac{1}{2}})$$

$$R_3(x) := \log(x) - x + \frac{x^2}{2}$$

$$\Delta(m) = \sum_{i=1}^k -m_i \left(u_i - \frac{1}{2} u_i^2 + R_3(u_i) \right) - \frac{1}{2} u_i + \frac{1}{4} u_i^2 - \frac{1}{2} R_3(u_i) + O(n^{\varepsilon - \frac{1}{2}})$$

$$\Delta(m) = \sum_{i=1}^k -m_i \left(u_i - \frac{1}{2} u_i^2 + R_3(u_i) \right) + \sum_{i=1}^k -\frac{1}{2} u_i + \frac{1}{4} u_i^2 - \frac{1}{2} R_3(u_i) + O(n^{\varepsilon - \frac{1}{2}})$$

Note that by definition $h_i = m_i^* u_i$ and $m_i = m_i^* (1 + u_i)$ so $m_i u_i = m_i^* (1 + u_i) u_i = h_i + m_i^* u_i^2$

Therefore since $\sum h_i = 0$ we get that

$$\Delta(m) = \sum_{i=1}^k -m_i^* u_i^2 + \frac{1}{2} m_i u_i^2 - m_i R_3(u_i) + \sum_{i=1}^k -\frac{1}{2} u_i + \frac{1}{4} u_i^2 - \frac{1}{2} R_3(u_i) + O(n^{\varepsilon - \frac{1}{2}})$$

$$\Delta(m) = \sum_{i=1}^k -\frac{1}{2} m_i^* u_i^2 + \frac{1}{2} h_i u_i^2 - m_i R_3(u_i) + \sum_{i=1}^k -\frac{1}{2} u_i + \frac{1}{4} u_i^2 - \frac{1}{2} R_3(u_i) + O(n^{\varepsilon - \frac{1}{2}})$$

Now use the fact that $u_i = \frac{h_i}{m_i^*}$ to get

$$\Delta(m) = \sum_{i=1}^k -\frac{1}{2} \frac{h_i^2}{m_i^*} + \frac{1}{2} \frac{h_i^3}{(m_i^*)^2} - m_i R_3(u_i) + \sum_{i=1}^k -\frac{1}{2} u_i + \frac{1}{4} u_i^2 - \frac{1}{2} R_3(u_i) + O(n^{\varepsilon - \frac{1}{2}})$$

$$\Delta(m) = \sum_{i=1}^k -\frac{1}{2} \frac{h_i^2}{m_i^*} + \frac{1}{2} \frac{h_i^3}{(m_i^*)^2} - m_i R_3(u_i) - \sum_{i=1}^k \frac{1}{2} R_3(u_i) + O(n^{\varepsilon - \frac{1}{2}})$$

Because $\sum u_i = \sum O(n^{\varepsilon - \frac{1}{2}}) = O(n^{\varepsilon - \frac{1}{2}})$ since the sum is finite and $\sum u_i^2$ gets even smaller, in fact it is $O(n^{2\varepsilon - 1})$ as n gets large.

Now $\frac{h_i^3}{(m_i^*)^2} = \frac{O(n^{\frac{3}{2} + 3\varepsilon})}{O(n^2)}$ where crucially for large enough n $(m_i^*)^2 \geq Cn^2$ for some positive C which allows us to say that its reciprocal is $O(n^{-2})$. It is here we use the fact that $\varepsilon < \frac{1}{6}$, as we now have

$$\Delta(m) = \sum_{i=1}^k -\frac{1}{2} \frac{h_i^2}{m_i^*} - m_i R_3(u_i) - \sum_{i=1}^k \frac{1}{2} R_3(u_i) + O(n^{3\varepsilon - \frac{1}{2}})$$

Now by Taylor's theorem, $R_3(u_i)$ from how it was defined is $\frac{1}{3} \xi_i^3$ for some $0 < \xi_i < u_i$, and since by definition $u_i = O(n^{\varepsilon - \frac{1}{2}})$ and $m_i = O(n)$ so $R_3(u_i) = O(n^{3\varepsilon - \frac{3}{2}})$ and thus so is $\sum_{i=1}^k \frac{1}{2} R_3(u_i)$, and also $m_i R_3(u_i) = O(n^{3\varepsilon - \frac{1}{2}})$ so at last we have

$$\Delta(m) = \sum_{i=1}^k -\frac{1}{2} \frac{h_i^2}{m_i^*} + O(n^{3\varepsilon - \frac{1}{2}})$$

The last step is

$$\Delta(m) = \sum_{i=1}^k -\frac{1}{2} \frac{h_i^2}{E_i + O(1)} + O(n^{3\varepsilon - \frac{1}{2}})$$

$$\Delta(m) = \sum_{i=1}^k -\frac{1}{2} \frac{h_i^2}{E_i} \left(1 + O\left(\frac{1}{n}\right)\right) + O(n^{3\varepsilon - \frac{1}{2}})$$

$$\Delta(m) = \sum_{i=1}^k -\frac{1}{2} \frac{h_i^2}{E_i} + \sum_{i=1}^k -\frac{1}{2} \frac{h_i^2}{E_i} O\left(\frac{1}{n}\right) + O(n^{3\varepsilon - \frac{1}{2}})$$

$$\Delta(m) = \sum_{i=1}^k -\frac{1}{2} \frac{h_i^2}{E_i} + O(n^{3\varepsilon - \frac{1}{2}})$$

Since that term was $O(n^{2\varepsilon - 1})$.

Now we are basically there because:

$$\Delta(m) = \log(P(m)) - \log(P(m^*))$$

Therefore

$$\log(P(m)) - \log(P(m^*)) = \frac{1}{2} |y|^2 + O(n^{3\varepsilon - \frac{1}{2}})$$

Undoing logs gives

$$\frac{P(m)}{P(m^*)} = e^{-\frac{1}{2}|y|^2} T$$

Where T goes to 1 as n goes to infinity for each fixed set of p values and regardless of m provided m is in R, so done with part 2 of the claim.

Lemma 2: There exists constants $c, C > 0$ such that $\sup_{m \notin R} p(m) \leq Cp(m^*)e^{-cn^{2\varepsilon}}$ for sufficiently large n.

Assume the probabilities are fixed and non-zero and the number of cells k is fixed.

Proof of lemma 2:

We first show the simple inequality that says $(1+t) \log(1+t) \geq t + \frac{t^2}{3}$ for all $t \in \left[-\frac{1}{2}, \frac{1}{2}\right]$.

Consider $(1+t) \log(1+t) - t - \frac{t^2}{3}$. The derivative of this is $\log(1+t) - \frac{2t}{3}$ and the second derivative is $\frac{1}{1+t} - \frac{2}{3}$ which in the range we care about is always non-negative. Notice that the derivative at 0 is 0 so we have a graph with positive curvature that flattens at $t=0$ where the value is 0, and thus the graph is non-negative in the range we care about so we finished the inequality.

Define a probability vector $q := (q_1, q_2, \dots, q_k), \sum q_r = 1$. Define $t_r := \frac{q_r}{p_r} - 1$.

Define $D(q||p) := \sum_{r=1}^k q_r \log\left(\frac{q_r}{p_r}\right) = \sum_{r=1}^k p_r(1+t_r) \log(1+t_r)$. Using the inequality above we have the bound $D(q||p) \geq \sum_{r=1}^k p_r \left(t_r + \frac{t_r^2}{3}\right)$ whenever all t_r 's are between $-\frac{1}{2}$ and $+\frac{1}{2}$ —we will handle the case where this is not true separately.

$$\begin{aligned} D(q||p) &\geq \sum_{r=1}^k p_r t_r + \frac{1}{3} \sum_{r=1}^k p_r t_r^2 = \sum_{r=1}^k q_r - p_r + \frac{1}{3} \sum_{r=1}^k p_r t_r^2 = \frac{1}{3} \sum_{r=1}^k p_r t_r^2 \\ &= \frac{1}{3} \sum_{r=1}^k \frac{(q_r - p_r)^2}{p_r} \end{aligned}$$

Given a table m set $q_r = \frac{m_r}{n}$, ie the observed probability. The standardized vector in components is given by $y_r = \sqrt{n} \frac{q_r - p_r}{\sqrt{p_r}} = \frac{O_r - E_r}{\sqrt{n p_r}}$. Recall that from this we derived $|y|^2 = \sum_{r=1}^k \frac{(O_r - E_r)^2}{E_r}$, and also it is the case that $|m - E|^2 = \sum_{r=1}^k (O_r - E_r)^2$. $|y|^2 = \sum_{r=1}^k \frac{(O_r - E_r)^2}{E_r} \geq \frac{1}{\max_r E_r} \sum_{r=1}^k (O_r - E_r)^2 = \frac{|m - E|^2}{\max_r E_r}$. By definition of R, if $m \notin R$, $|y|^2 \geq \frac{n^{1+2\varepsilon}}{\max_r E_r} = \frac{n^{2\varepsilon}}{p_{\max}}$. Here $p_{\max}$ is the largest probability.

In the case where all t_r 's are between $-\frac{1}{2}$ and $+\frac{1}{2}$ we had $D(q||p) \geq \frac{1}{3} \sum_{r=1}^k \frac{(q_r - p_r)^2}{p_r} = \frac{1}{3n} \sum_{r=1}^k \frac{(O_r - E_r)^2}{E_r} = \frac{|y|^2}{3n}$ so $nD(q||p) \geq \frac{y^2}{3}$ so by the previous deduction $nD(q||p) \geq \frac{n^{2\varepsilon}}{3p_{\max}}$.

If for some r, $|t_r| > \frac{1}{2}$ then the set of possible q vectors (including cases where some are 0) satisfying $|t_r| \geq \frac{1}{2}$ is bounded and since it contains all its boundary points it is closed, and trivially it does not contain p. The function $D(q||p)$ for fixed p is continuous on K (we say $0 \log(0) = 0$ for the case when one of the q's is 0 as that is the limit of $x \log(x)$ as x goes to 0), it must be the case (by the technical results

document, extreme value theorem) that $D(q|p)$ attains a minimum value on K d_0 somewhere in K . Note that $d_0 > 0$: The proof of this is as follows.

For all positive u , $\log(u) \leq u - 1$ because those curves are tangent at $u=1$ and $\log$ has negative curvature everywhere (differentiate it twice). Therefore $\log\left(\frac{p_r}{q_r}\right) \leq \frac{p_r}{q_r} - 1$ so $q_r \log\left(\frac{p_r}{q_r}\right) \leq p_r - q_r$, therefore $q_r \log\left(\frac{q_r}{p_r}\right) \geq q_r - p_r$, so $\sum_{r=1}^k q_r \log\left(\frac{q_r}{p_r}\right) \geq \sum_{r=1}^k q_r - p_r = 0$.

Note that $\log(u) = u - 1$ only happens when $u = 1$ so if p and q differ for any r then that term in the sum becomes strictly positive and the rest are non-negative so the inequality becomes strict, hence the inequality is strict whenever q and p are not the same, so certainly everywhere in K .

Therefore everywhere in K , $nD(q|p) \geq nd_0$.

Note that as explained earlier the number of ways to get a certain table by doing a bunch of trials is exactly $\frac{n!}{\prod m_r!}$, and $\prod p_r^{m_r} = \exp(n \sum q_r \log(p_r))$ where $\exp(x)$ means e^x . Now define another function $H(q) := -\sum q_r \log(q_r)$ with $0\log(0)$ defined as 0.

Now $\frac{n!}{\prod m_r!}$ is the $\prod_{r=1}^k m_r^{m_r}$ coefficient in the expansion of $n^n = (m_1 + m_2 + \dots + m_k)^n$ so we have the inequality $\frac{n!}{\prod m_r!} \leq \frac{n^n}{\prod_{r=1}^k m_r^{m_r}}$ so $p(m) = \frac{n!}{\prod m_r!} \prod p_r^{m_r} \leq \prod p_r^{m_r} \frac{n^n}{\prod_{r=1}^k m_r^{m_r}}$.

Taking logs gives $\log(p(m)) \leq n \log(n) - \sum m_r \log(m_r) + \sum m_r \log(p_r)$.

$$\begin{aligned} \log(p(m)) &\leq n \log(n) - \sum n q_r \log(q_r) - \sum n q_r \log(n) + \sum m_r \log(p_r) \\ \log(p(m)) &\leq - \sum n q_r \log(q_r) + \sum n q_r \log(p_r) \\ \log(p(m)) &\leq n \sum (q_r) \log\left(\frac{p_r}{q_r}\right) = -nD(q|p) \end{aligned}$$

Thus

$$p(m) \leq e^{-nD(q|p)} \leq e^{-nd_0}$$

Now we apply lemma 1 and take logs of both sides of the formula for probabilities of particular unconstrained tables.

$$\log(p(m)) = \log(n!) - \sum_{r=1}^k \log(m_r!) + \sum_{r=1}^k m_r \log(p_r)$$

$$\log(p(m)) = n \log(n) - n + \frac{1}{2} \log(2n\pi) + R_n - \sum_{r=1}^k m_r \log(m_r) - m_r + \frac{1}{2} \log(2m_r\pi) + R_{m_r} + \sum_{r=1}^k m_r \log(p_r)$$

Where $0 < R_n < \frac{1}{12n}$ and $0 < R_{m_r} < \frac{1}{12m_r} \leq \frac{1}{12np_{min}}$, therefore

$$\log(p(m)) = n \log(n) - n + \frac{1}{2} \log(2n\pi) - \sum_{r=1}^k \left[m_r \log(m_r) - m_r + \frac{1}{2} \log(2m_r\pi) \right] + \sum_{r=1}^k m_r \log(p_r) + R$$

Where R goes to 0 as n goes to infinity.

Simplifying,

$$\log(p(m)) = n \log(n) + \frac{1}{2} \log(2n\pi) - \sum_{r=1}^k \left[m_r \log(m_r) + \frac{1}{2} \log(2m_r\pi) \right] + \sum_{r=1}^k m_r \log(p_r) + R$$

Because n is the sum of the m_r 's

$$\log(p(m)) = n \log(n) + \frac{1}{2} \log(2n\pi) - \sum_{r=1}^k [m_r \log(m_r)] - \frac{1}{2} \sum_{r=1}^k [\log(2m_r\pi)] + \sum_{r=1}^k m_r \log(p_r) + R$$

$$\log(p(m)) = \frac{1}{2} \log(2n\pi) - \sum_{r=1}^k [m_r [\log(q_r)]] - \frac{1}{2} \sum_{r=1}^k [\log(2m_r\pi)] + \sum_{r=1}^k m_r \log(p_r) + R$$

$$\log(p(m)) = \frac{1}{2} \log(2n\pi) - \sum_{r=1}^k [m_r [\log\left(\frac{q_r}{p_r}\right)]] - \frac{1}{2} \sum_{r=1}^k [\log(2m_r\pi)] + R$$

$$\log(p(m)) = \frac{1}{2} \log(2n\pi) - n \sum_{r=1}^k \left[q_r \left[\log\left(\frac{q_r}{p_r}\right) \right] \right] - \frac{1}{2} \sum_{r=1}^k [\log(2m_r\pi)] + R$$

$$\log(p(m)) = \frac{1}{2} \log(2n\pi) - nD(q \mid p) - \frac{1}{2} \sum_{r=1}^k [\log(2m_r\pi)] + R$$

$$\log(p(m)) = \frac{1-k}{2} \log(n) - nD(q \mid p) - \frac{1}{2} \sum_{r=1}^k [\log(q_r)] - \frac{k+1}{2} \log(2\pi) + R$$

Note that $\frac{k+1}{2} \log(2\pi)$ is a fixed constant here.

Note that $|m^* - E| \leq \sqrt{k}$ from earlier, therefore $|q^* - p| \leq \frac{\sqrt{k}}{n}$ and therefore $D(q \mid p)$.

By the technical results document (taylor's theorem),

$$\begin{aligned} \log\left(\frac{q_r}{p_r}\right) &= \log\left(1 + \frac{q_r - p_r}{p_r}\right) = \frac{q_r - p_r}{p_r} - \frac{1}{2} \left(\frac{q_r - p_r}{p_r}\right)^2 + O\left(\left(\frac{q_r - p_r}{p_r}\right)^3\right). \text{ But then } q_r \log\left(\frac{q_r}{p_r}\right) = \\ &= (q_r) \left[\frac{q_r - p_r}{p_r} - \frac{1}{2} \left(\frac{q_r - p_r}{p_r}\right)^2 + O\left(\left(\frac{q_r - p_r}{p_r}\right)^3\right) \right]. \text{ Now set } \delta_r := q_r - p_r, \text{ and we know from what we did a few} \\ &\text{lines ago that } \delta_r = O\left(\frac{1}{n}\right). \text{ } q_r \log\left(\frac{q_r}{p_r}\right) = (p_r + \delta_r) \left[\frac{\delta_r}{p_r} - \frac{1}{2} \left(\frac{\delta_r}{p_r}\right)^2 + O\left(\left(\frac{\delta_r}{p_r}\right)^3\right) \right] \\ &= \delta_r - \frac{1}{2} \frac{\delta_r^2}{p_r} + \frac{\delta_r^2}{p_r} + O(\delta_r^3) = \delta_r + \frac{1}{2} \frac{\delta_r^2}{p_r} + O(\delta_r^3) \end{aligned}$$

Note however that when we sum over r to get $D(q^* \mid p)$, the δ_r term vanishes since those add to 0, and since $\delta_r = O\left(\frac{1}{n}\right)$, $D(q^* \mid p) = O\left(\frac{1}{n^2}\right)$ and therefore $nD(q^* \mid p) = O\left(\frac{1}{n}\right)$. Since $q_r^* - p_r$, the log terms stay bounded.

Therefore we have

$$\log(p(m)) = \frac{1-k}{2} \log(n) - nD(q \mid p) - \frac{1}{2} \sum_{r=1}^k [\log(q_r)] + O(1)$$

From earlier, and

$$\log(p(m^*)) = \frac{1-k}{2} \log(n) - \frac{1}{2} \sum_{r=1}^k [\log(q_r^*)] + O(1)$$

Since the terms $+\frac{1}{2} \sum_{r=1}^k [\log(q_r^*)] + O(1)$ stay bounded, we reduce to

$$\log(p(m^*)) = \frac{1-k}{2} \log(n) + O(1)$$

Therefore

$$\log\left(\frac{p(m)}{p(m^*)}\right) = -nD(q|p) - \frac{1}{2} \sum_{r=1}^k [\log(q_r)] + O(1)$$

Therefore we just need to bound $\sum_{r=1}^k [\log(q_r)]$. Outside of K, this is fine, as each q is bounded.

Therefore we have that

$$\log\left(\frac{p(m)}{p(m^*)}\right) = -nD(q|p) + O(1)$$

Therefore, if n is large enough, $\frac{p(m)}{p(m^*)} \leq C e^{-nD(q|p)}$ outside K for some C provided n is large enough.

Since

$$\log(p(m^*)) = \frac{1-k}{2} \log(n) + O(1)$$

It means that for n large enough, there exists positive constants A and B such that

$$A n^{\frac{1-k}{2}} < p(m^*) < B n^{\frac{1-k}{2}}$$

Therefore, we have the bound (for some positive constants A and B).

Therefore, because outside K, $nD(q|p) \geq \frac{n^{2\varepsilon}}{3p_{max}}$, $\frac{p(m)}{p(m^*)} \leq C e^{-cn^{2\varepsilon}}$ so the lemma is proved outside K.

Inside K, $p(m) \leq e^{-d_0 n}$ so $\frac{p(m)}{p(m^*)} \leq C n^{\frac{k-1}{2}} e^{-d_0 n} \leq e^{-\frac{d_0}{2} n} \leq e^{-\frac{d_0}{3} n^{2\varepsilon}}$ for some C (because $\varepsilon < \frac{1}{2}$) and any large enough n. The second inequality is because $C n^{\frac{k-1}{2}}$ eventually is beaten by $e^{\frac{d_0}{2} n}$ (take the ratio and apply lhopital at least $\frac{k-1}{2}$ times as n goes to infinity to see this) This proves the lemma in all cases.

Proof of part 3 of the claim:

Pick A from the statement of the lemma to be $\sqrt{c p_{min}}$ where c is as in the lemma 2 statement and p_{min} is the smallest probability in the table.

Define R^C as the outside of R. Pick $m_0 \in S \cap R_2$, then $P(m \notin R | m \in S) = \frac{\sum_{m \in S \cap R^C} p(m)}{\sum_{m \in S} p(m)} \leq \frac{\sum_{m \in S \cap R^C} p(m)}{p(m_0)}$.

Let |S| be the number of lattice points in S. Then $\sum_{m \in S \cap R^C} p(m) \leq |S| \sup_{m \in R} p(m)$. By lemma 2, if n is

large enough, $\sup_{m \in R} p(m) \leq |S| C p(m^*) e^{-cn^{2\varepsilon}}$. By part 2, if n is large enough, since $m_0 \in R$ we can force

$p(m_0) \geq \frac{1}{2} p(m^*) e^{-\frac{1}{2} |y_0|^2}$ where y_0 is the vector m_0 moved to the standardized world.

Now since the standardization is $y_i = \frac{m_i - E_i}{\sqrt{np_i}}$ we have

$$|y_0|^2 = \sum_{i=1}^k \frac{((m_0)_i - E_i)^2}{E_i} \leq \frac{1}{np_{\min}} \sum_{i=1}^k ((m_0)_i - E_i)^2 = \frac{|m_0 - E|^2}{np_{\min}}$$

Since $m_0 \in R_2$, $|y_0|^2 \leq \frac{1}{np_{\min}} cp_{\min} n^{1+2\varepsilon}$ by how R_2 is defined, and this is just $|y_0|^2 \leq cp_{\min} \frac{n^{2\varepsilon}}{p_{\min}}$.

Therefore we have that $p(m_0) \geq \frac{1}{2}p(m^*)e^{-\frac{1}{2}cn^{2\varepsilon}}$. Going back to the original inequality and putting in all our results, we get the following:

$$P(m \notin R | m \in S) \leq \frac{\sum_{m \in S \cap R^c} p(m)}{p(m_0)} \leq |S| \frac{Cp(m^*)e^{-cn^{2\varepsilon}}}{\frac{1}{2}p(m^*)e^{-\frac{1}{2}cn^{2\varepsilon}}} = 2C|S|e^{-\frac{1}{2}cn^{2\varepsilon}}$$

Now there are $n+1$ possibilities for each cell and k cells so we get the trivial upper bound which says, therefore $|S| \leq (n+1)^k$

$$P(m \notin R | m \in S) \leq 2C(n+1)^k e^{-\frac{1}{2}cn^{2\varepsilon}}$$

Again exponentials beat polynomials so the above goes to 0, completing the proof of the proposition. So done.

As for why we combine cells with expected values under 5, it's just a rule of thumb, not really a rigorous statement with anything special about 5, at least as far as I know. Essentially the reason is that for the central limit theorem to work, it has to be possible for us to get tails in both directions. If the expected value was, say 1, then the distribution would not possibly be able to be symmetrical without going into the negative values. Also, if p is small (which is usually when we have small expected values), the cell values become well approximated by a poisson distribution, and by the central limit theorem, the poisson distribution is better approximated by a normal the larger the mean is. As for why 5? You can see from this picture that is about when the poisson distribution starts to look like a normal curve.




Image: Po(3, 4, 5, 6) on the same graph, demonstrating that around 5 is when it starts to look like a normal distribution.

The formula $(n-1) \frac{S^2}{\sigma^2} \sim \chi_{n-1}^2$

This follows from the definition of S^2 , σ^2 and χ_{n-1}^2 . The result is clear if you were doing this with a $N(0,1)$ variable, then observe that $\frac{S^2}{\sigma^2}$ does not change when we scale or move the variable since s scales with σ and neither change by shifting. Here is why it works for an $N(0,1)$ variable: If I have a bunch of independently distributed random variables, then we can use the same trick of using rotational symmetry of the normal distribution.

Let's define

$$Z_i = \frac{X_i - \mu}{\sigma} \sim N(0, 1)$$

so $X_i = \mu + \sigma Z_i$.

We have

$$\sum_{i=1}^n (X_i - \bar{X})^2 = \sigma^2 \sum_{i=1}^n (Z_i - \bar{Z})^2$$

where $\bar{Z} = \frac{1}{n} \sum_{i=1}^n Z_i$.

Therefore,

$$(n-1) \frac{S^2}{\sigma^2} = \sum_{i=1}^n (Z_i - \bar{Z})^2$$

- The vector $(Z_1, \dots, Z_n)$ is a random point in $\mathbb{R}^n$, distributed as a standard normal vector.
- The **mean** $\bar{Z}$ is the projection of this vector onto the direction $(1, 1, \dots, 1)$.
- The quantity $\sum_{i=1}^n (Z_i - \bar{Z})^2$ is the squared length of the component **orthogonal** to $(1, 1, \dots, 1)$.

By rotational symmetry:

- The "length squared" of the orthogonal component is distributed as χ_{n-1}^2 .

This is because a standard normal confined to the plane such that the sample mean is what it is is just a rotation of a normal with one fewer variable in a normal plane, and in that case chi squared is clearly the distance squared from the origin of that multivariate normal. Conditioning is allowed as it is a continuous distribution and issues came from using a discrete distribution.

t test for difference between means

Note: I'm not really sure how the degrees of freedom relate to the dimension test, so I will merely prove the mathematical properties required.

Note that a t_{n-1} distribution is given by $\frac{\bar{X} - \mu}{\frac{S}{\sqrt{n}}}$ in a sample of size n normally distributed variables, where scaling them will affect the numerator and the denominator in the same way and shifting them will not affect the numerator. In the two-sample t-test, we define the following, assuming X and Y have equal variance:

$$T = \frac{\bar{X} - \bar{Y} - (\mu_X - \mu_Y)}{\sqrt{\frac{1}{n_X} + \frac{1}{n_Y}} S_p}, \quad S_p^2 = \frac{\sum_{i=1}^{n_X} (X_i - \bar{X})^2 + \sum_{j=1}^{n_Y} (Y_j - \bar{Y})^2}{n_X + n_Y - 2},$$

Note that S_p^2 is an unbiased estimate for the variance since it is a linear combination of unbiased estimates for $\text{Var}(X)$ and $\text{Var}(Y)$. We ultimately need to show that T has the same distribution as a $t_{n_X+n_Y-2}$ variable. Note that we have the following:

$$Z := \frac{\bar{X} - \bar{Y} - (\mu_X - \mu_Y)}{\sigma \sqrt{\frac{1}{n_X} + \frac{1}{n_Y}}} \text{ is } N(0, 1)$$

The reason is because the variance of $n_x \bar{X}$, which is the sum of the elements of X, by additivity of variances is the sum of all the variances of the elements of x which is $n_x \sigma^2$. Therefore, since shifting does not affect the variance this is also the variance of $n_x(\bar{X} - \mu)$. Since multiplying a variable by the square of a constant multiplies the variance by that constant, we get that the variance of $\bar{X} - \mu$ must be $\frac{\sigma^2}{n_x}$. Similarly, we have that the variance of $\bar{Y} - \mu$ is $\frac{\sigma^2}{n_y}$. Therefore, we have that the expression above is standardized correctly since the variance of the numerator is $\frac{\sigma^2}{n_x} + \frac{\sigma^2}{n_y}$. We also have the following:

$$U := \frac{(n_X + n_Y - 2) S_p^2}{\sigma^2} \text{ is } \chi^2_{n_X+n_Y-2}$$

For the same reason as the similar formula we proved above: We simply have to constrain the means of X and Y which reduces the dimension by 2.

Now we have that

$$T = \frac{\frac{Z}{\sigma}}{\sqrt{\frac{U}{n_X+n_Y-2}}} \text{ by definition of T. However, this is equal to } \frac{\frac{N(0,1)}{\sqrt{\frac{\chi^2_{n_X+n_Y-2}}{n_X+n_Y-2}}}}{\sqrt{\frac{\chi^2_{n_X+n_Y-2}}{n_X+n_Y-2}}}$$

variable in the traditional sense using the relation between S^2 and χ^2 .

Other stuff about t and f distributions do not require mathematics since the tables are based on numerical data with no other implicit assumptions other than invariance under shifting and scaling.